

Can We Automate Scientific Reviewing?

Weizhe Yuan

*New York University,
New York, NY 10003*

wy885@nyu.edu

Pengfei Liu

*Carnegie Mellon University,
Pittsburgh, PA 15213*

pliu3@cs.cmu.edu

Graham Neubig

*Carnegie Mellon University,
Pittsburgh, PA 15213*

gneubig@cs.cmu.edu

Abstract

The rapid development of science and technology has been accompanied by an exponential growth in peer-reviewed scientific publications. At the same time, the review of each paper is a laborious process that must be carried out by subject matter experts. Thus, providing high-quality reviews of this growing number of papers is a significant challenge. In this work, we ask the question “*can we automate scientific reviewing?*”, discussing the possibility of using natural language processing (NLP) models to generate peer reviews for scientific papers. Because it is non-trivial to define what a “good” review is in the first place, we first discuss possible evaluation metrics that could be used to judge success in this task. We then focus on the machine learning domain and collect a dataset of papers in the domain, annotate them with different aspects of content covered in each review, and train targeted summarization models that take in papers as input and generate reviews as output. Comprehensive experimental results on the test set show that while system-generated reviews are *comprehensive*, touching upon more aspects of the paper than human-written reviews, the generated texts are *less constructive* and *less factual* than human-written reviews for all aspects except the explanation of the core ideas of the papers, which are largely factually correct. Given these results, we pose *eight* challenges in the pursuit of a good review generation system together with potential solutions, which, hopefully, will inspire more future research in this direction.

We make relevant resource publicly available for use by future research: <https://github.com/neulab/ReviewAdvisor>. In addition, while our conclusion is that the technology is not yet ready for use in high-stakes review settings we provide a system demo, *ReviewAdvisor* (<http://review.nlpedia.ai/>), showing the current capabilities and failings of state-of-the-art NLP models at this task (see demo screenshot in A.2). A review of this paper written by the system proposed in this paper can be found in A.1.

1. Introduction

The number of published papers is growing exponentially (Tabah, 1999; De Bellis, 2009; Bornmann & Mutz, 2015). While this may be positively viewed as indicating acceleration of scientific progress, it also poses great challenges for researchers, both in reading and synthesizing the relevant literature for one’s own benefit, and for performing *peer review* of papers to vet their correctness and merit. With respect to the former, a large body of existing work explores automatic summarization of a paper or a set of papers for automatic survey

generation (Mohammad et al., 2009; Jha et al., 2013, 2015b, 2015a; Yasunaga et al., 2019b; Cohan et al., 2018b; Xing et al., 2020). However, despite its importance, peer review is both laborious and error-prone, and automatic systems to aid in the peer review process remain relatively underexplored. Bartoli et al. (2016) investigated the feasibility of generating reviews by surface-level term replacement and sentence reordering, and Wang et al. (2020; contemporaneously and independently) propose a two-stage information extraction and summarization pipeline to generate paper reviews. However, both do not extensively evaluate the quality or features of the generated review text.

In this work, we are concerned with providing at least a preliminary answer to the ambitious over-arching question: *can we automate scientific reviewing?* Given the complexity of understanding and assessing the merit of scientific contributions, we do not expect an automated system to be able to match a well-qualified and meticulous human reviewer at this task any time soon. However, some degree of review automation may allow authors to iteratively refine and improve their papers, point out parts of the paper that reviewers may have missed in their assessments, or provide a reference to junior reviewers who are just learning the ropes of the reviewing process. Towards this goal, we examine two concrete research questions, the answers to which are prerequisites to building a functioning review assistant:

Q1: What are the desiderata of a good automatic reviewing system, and how can we quantify them for evaluation? Before developing an automatic review system, we first must quantify what constitutes a good review in the first place. The challenge of answering this question is that a review commonly involves both objective aspects (e.g. “lack of details necessary to replicate the experimental protocol”) and subjective aspects (e.g. “lack of potential impact”). Due to this subjectivity, defining a “good” review is itself somewhat subjective.

As a step towards tackling this challenge, we argue that it is possible to view review generation as a task of *aspect-based scientific paper summarization* but with some important differences, where the reviewer not only tries to summarize the core idea of a paper, but also gives *critical* comments based on his/her understanding of the *whole literature* from diverse aspects (e.g. novelty or potential impact). We evaluate review quality from multiple perspectives, in which we claim a good review not only should make a good summary (§2.2) of a paper but also consist of factually correct (§2.2), critical (§2.2), and fair comments (§5.3) from diverse aspects (§2.2), together with informative evidence (§2.2).

To operationalize these concepts, we build a dataset of reviews, named **ASAP-Review**¹ from the machine learning domain, and make fine-grained annotations of aspect information for each review, which provides the possibility for a richer evaluation of generated reviews.

Q2: Using state-of-the-art NLP models, to what extent can we realize these desiderata? We provide an initial answer to this question by using the aforementioned dataset to train state-of-the-art summarization models to generate reviews from scientific papers, and evaluate the output according to our evaluation metrics described above. We propose different architectural designs for this model, which we dub *ReviewAdvisor* (§4), and comprehensively evaluate them, interpreting their relative advantages.

Main observations and conclusions:

1. ASpect-enhAnced Peer Review dataset

Finally, we highlight the observations and conclusions we find from our experiments on review generation (for machine learning papers).

(1) *What are review generation systems (not) good at?* Most importantly, we find the constructed automatic review system *generates non-factual statements* regarding many aspects of the paper assessment, which is a serious flaw in a high-stakes setting such as reviewing. However, there are some bright points as well. For example, it *can often precisely summarize the core idea* of the input paper, which can be either used as a draft for human reviewers or help them (or general readers) quickly understand the main idea of the paper to be reviewed (or pre-print papers). It can also generate reviews that *cover more aspects* of the paper’s quality than those created by humans, and provide evidence sentences from the paper. These could potentially provide a preliminary template for reviewers and help them more comprehensively identify salient information in making their assessment.

(2) *Will the system generate biased reviews?* Probably. We present methods to identify and quantify disparities (potential biases) in reviews (§5.3), and find that both human and automatic reviewers exhibit varying degrees of disparity. (i) regarding native vs. non-native English speakers: papers of native English speakers tend to obtain higher scores on “Clarity” from human reviewers than non-native English ones,² but the automatic review generators narrow this gap. Additionally, system reviewers are harsher than human reviewers when commenting regarding the paper’s “Originality” for non-native English speakers. (ii) regarding anonymous vs. non-anonymous submissions: both human reviewers and system reviewers favor non-anonymous papers, which have been posted on non-blind preprint servers such as arXiv³ before the review period, more than anonymous papers in all aspects. This finding is interesting because the proposed model does not explicitly reference arXiv, indicating underlying differences in the textual content of posted vs. unposted papers is responsible for the differences in system reviews, and perhaps human reviews as well.

Based on above mentioned issues, we claim that **a review generation system can not replace human reviewers at this time, instead, but it may, sooner or later, be helpful as part of a machine-assisted human review process**. Our research also suggests what may be next steps in pursuing better methods for automatic review generation or assistance, and we summarize *eight challenges* that can be explored in §7.2.

2. What Makes a Good Peer Review?

Although peer review has been adopted by most journals and conferences to identify important and relevant research, its effectiveness is being continuously questioned (Smith, 2006; Langford & Guzdial, 2015; Tomkins et al., 2017; Gao et al., 2019; Rogers & Augenstein, 2020).

As concluded by Jefferson et al.’s (2002b): “*Until we have properly defined the objectives of peer-review, it will remain almost impossible to assess or improve its effectiveness.*” Therefore we first discuss the possible objectives of peer review.

-
- 2. Whether this actually qualifies as “bias” is perhaps arguable. Papers written by native English speakers may be more clear due to lack of confusing grammatical errors, but the paper may actually be perfectly clear but give the impression of not being clear because of grammatical errors.
 - 3. <https://arxiv.org/>

2.1 Peer Review for Scientific Research

A research paper is commonly first reviewed by several reviewers who usually assign one or several *scores* and give detailed comments. The comments, and sometimes scores, cover diverse *aspects* of the paper such as “clarity” or “potential impact”, detailed in §3.2.1). These aspects are often directly mentioned in review forms of scientific conferences or journals.⁴ Then a senior reviewer will often make a *final decision* (i.e., “reject” or “accept”) and provide comments summarizing the decision (i.e., a *meta-review*).

After going through many review guidelines⁵ and resources about how to write a good review⁶ we summarize *some* of the most frequently mentioned desiderata below:

1. **Decisiveness:** A good review should take a clear stance, selecting high-quality submissions for publication and suggesting others not be accepted (Jefferson et al., 2002a; Smith, 2006).
2. **Comprehensiveness:** A good review should be well-organized, typically starting with a brief summary of the paper’s contributions, then following with opinions gauging the quality of a paper from different aspects. Many review forms explicitly require evaluation of different aspects to encourage comprehensiveness.
3. **Justification:** A good review should provide specific reasons for its assessment, particularly whenever it states that the paper is lacking in some aspect. This justification also makes the review more constructive (another oft-cited desideratum of reviews), as these justifications provide hints about how the authors could improve problematic aspects in the paper (Xiong & Litman, 2011).
4. **Accuracy:** A review should be factually correct, with the statements contained therein not being demonstrably false.
5. **Kindness:** A good review should be kind and polite in language use.

Based on above desiderata, we make a first step towards evaluation of reviews for scientific papers and characterize a “good” review from multiple perspectives.

2.2 Multi-Perspective Evaluation

Given input paper D and meta-review R^m , our goal is to evaluate the quality of review R , which can be either manually or automatically generated. We also introduce a function $\text{DEC}(D) \in \{1, -1\}$ that indicates the final decision of a given paper reached by the meta-review: “accept” or “reject”. Further, $\text{REC}(R) \in \{1, 0, -1\}$ represents the acceptance recommendation of a particular review: “accept,” “neutral,” or “reject” (see Appendix A.3 for details).

-
4. For example, one example from the Conference of the Association for Computational Linguistics (ACL) can be found at: https://acl2018.org/downloads/acl_2018_review_form.html
 5. <https://icml.cc/Conferences/2020/ReviewerGuidelines> <https://NeurIPS.cc/Conferences/2020/PaperInformation/ReviewerGuidelines>, <https://iclr.cc/Conferences/2021/ReviewerGuide>
 6. https://players.brightcove.net/3806881048001/rFXiCa5uY_default/index.html?videoId=4518165477001, <https://soundcloud.com/nlp-highlights/77-on-writing-quality-peer-reviews-with-noah-a-smith>, <https://www.aclweb.org/anthology/2020.acl-tutorials.4.pdf>, <https://2020.emnlp.org/blog/2020-05-17-write-good-reviews>

Below, we discuss evaluation metrics that can be used to approximate the desiderata of reviews described in the previous section (summarized in Tab. 1).

Desiderata	Decisive.	Justification	Comprehen.		Accuracy		Others	
Metrics	RAcc	INFO	ACov	AREC	SAcc	ACon	ROUGE	BERTScore
Range	[−1, 1]	[0, 1]	[0, 1]	[0, 1]	[0, 1]	[0, 1]	[0, 1]	[0, 1]
Automated	No	No	Yes	Yes	No	No	Yes	Yes

Table 1: Evaluation metrics from different perspectives. “Range” represents the range value of each metric. “Automated” denotes if metrics can be obtained automatically.

2.2.1 D1: DECISIVENESS

First, we tackle the *decisiveness*, as well as accuracy of the review’s recommendation, through **Recommendation Accuracy (RAcc)**. Here we use the final decision regarding a paper and measure whether the acceptance implied by the review R is consistent with the actual accept/reject decision of the reviewed paper D . It is calculated as:

$$\text{RAcc}(R) = \text{DEC}(D) \times \text{REC}(R) \quad (1)$$

The highest score can only be achieved by making a correct acceptance recommendation (i.e., both $\text{DEC}(D)$ and $\text{REC}(R)$ take a value of 1 or -1). A higher score indicates that the review more decisively and accurately makes an acceptance recommendation.

2.2.2 D2: COMPREHENSIVENESS

A comprehensive review should touch on the quality of different aspects of the paper, which we measure using a metric dubbed **Aspect Coverage (ACov)**. Specifically, given a review R , aspect coverage measures how many aspects (e.g. clarity) in a predefined aspect typology (in our case, §3.2.1) have been covered by R .

In addition, we propose another metric **Aspect Recall (AREC)**, which explicitly takes the meta-review R^m into account. Because the meta-review is an authoritative summary of all the reviews for a paper, it provides an approximation of which aspects, and with which sentiment polarity, should be covered in a review. Aspect recall counts how many aspects in meta-review R^m are covered by general review R , with higher aspect recall indicating better agreement with the meta-review.⁷

2.2.3 D3: JUSTIFICATION

As defined in §2.1, a good peer review should provide hints about how the author could improve problematic aspects. For example, when reviewers comment: “this paper lacks important references”, they should also list these relevant works. To satisfy this justification

7. Notably, this metric potentially biases towards high scores for reviews that were considered in the writing of the meta-review. Therefore, higher aspect recall is not the only goal, and should be taken together with other evaluation metrics.

desideratum, we define a metric called **Informativeness (Info)** to quantify how many negative comments⁸ are accompanied by corresponding evidence.

First, let $n_{na}(R)$ denote the number of aspects in R with negative sentiment polarity. $n_{nae}(R)$ denotes the number of aspects with negative sentiment polarity that are supported by evidence. The judgement of supporting evidence is conducted manually (details in Appendix A.3). INFO is calculated as:

$$\text{Info}(R) = \frac{n_{nae}(R)}{n_{na}(R)} \quad (2)$$

And we set it to be 1 when there are no negative aspects mentioned in a review.

2.2.4 D4: ACCURACY

We use two measures to evaluate the accuracy of assessments. First, we use **Summary Accuracy (SAcc)** to measure how well a review summarizes contributions of a paper. It takes value of 0, 0.5, or 1, which evaluates the summary part of the review as incorrect/absent, partially correct, and correct. The correctness judgement is performed manually, with details listed in Appendix A.3.

Info implicitly requires that negative aspects should be supported with evidence, ignoring the quality of this evidence. However, to truly help to improve the quality of a paper, the evidence for negative aspects should be factual as well. Here we propose **Aspect-level Constructiveness (ACon)**, the percentage of the supporting statements $n_{nae}(R)$ that are judged as valid support by human annotators. If $n_{nae}(R)$ is 0, we set its **ACon** as 1. This metric will implicitly favor reviews that do not provide enough evidence for negative aspects. However, in this case, the **Info** of those reviews will be rather low. The details of evaluating “validity” are also described in Appendix A.3.

2.2.5 D5: KINDNESS

While kindness is very important in maintaining a positive research community, accurately measuring it computationally in a nuanced setting such as peer review is non-trivial. Thus, we leave the capturing of kindness in evaluation to future work.

2.2.6 SIMILARITY TO HUMAN REVIEWS

For automatically generated reviews, we also use **Semantic Equivalence** metrics to measure the similarity between generated reviews and reference reviews. The intuition is that while human reviewers are certainly not perfect, knowing how close our generated reviews are to existing human experts may be informative. Here, we investigate two specific metrics: **ROUGE** (Lin & Hovy, 2003) and **BERTScore** (Zhang et al., 2019). The former measures the surface-level word match while the latter measures the distance in embedding space. Notably, for each source input, there are multiple reference reviews. When aggregating **ROUGE** and **BERTScore**, we take the maximum instead of average since it is not necessary for generated reviews to be close to all references.

8. We only consider whether the reviewer has provided enough evidence for negative opinions since we find that most human reviewers rarely provide evidence for their positive comments.

3. Dataset

Next, in this section we introduce how we construct a review dataset with more fine-grained metadata, which can be used for system training and the multiple perspective evaluation of reviews.

3.1 Data Collection

The advent of the Open Peer Review system⁹ makes it possible to access review data for analysis or model training/testing. One previous work (Kang et al., 2018) attempts to collect reviews from several prestigious publication venues including the Conference of the Association of Computational Linguistics (ACL) and the International Conference on Learning Representations (ICLR). However, there were not nearly as many reviews accumulated in OpenReview at that time¹⁰ and other private reviews only accounted for a few hundred. Therefore we decided to collect our own dataset **Aspect-enhanced Peer Review (ASAP-Review)**.

We crawled ICLR papers from 2017-2020 through OpenReview¹¹ and NeurIPS papers from 2016-2019 through NeurIPS Proceedings.¹² For each paper’s review, we keep as much metadata information as possible. Specifically, for each paper, we include following metadata information that we can obtain from the review web page:

- *Reference reviews*, which are written by a committee member.
- *Meta reviews*, which are commonly written by an area chair (senior committee member).
- *Decision*, which denotes a paper’s final “accept” or “reject” decision.
- *Other* information like url, title, author, etc.

We used Allenai Science-parse¹³ to parse the pdf of each paper and keep the structured textual information (e.g., titles, authors, section content and references). The basic statistics of our **ASAP-Review** dataset is shown in Tab. 2¹⁴.

	Accept	Reject	Total	Avg. Full Text Len.	Avg. Review Len.	# of Reviews
ICLR	1,859	3,333	5,192	7,398	445	15,728
NeurIPS	3,685	0	3,685	5,916	411	12,391
Both	5,544	3,333	8,877	6,782	430	28,119

Table 2: Basic statistics of **ASAP-Review** dataset. Note that NeurIPS only provide reviews for accepted papers to the public.

9. <https://openreview.net/>

10. During that time, there are no reviews of ICLR from 2018 to 2020 nor reviews of NeurIPS from 2018 to 2019.

11. <https://openreview.net>

12. <http://papers.NeurIPS.cc>

13. <https://github.com/allenai/science-parse>

14. Although NeurIPS does not provide reviews of rejected papers, we assume that ICLR’s reviews and NeurIPS’s reviews are similar. We also agree that designing a domain-dependent (e.g., venue-dependent) peer review system is a promising future direction.

3.2 Aspect-enhanced Review Dataset

Although reviews exhibit internal structure, for example, as shown in Fig. 3, reviews commonly start with a paper summary, followed by different aspects of opinions, together with evidence. In practice, this useful structural information cannot be obtained directly. Considering that fine-grained information about the various aspects touched on by the review plays an essential role in review evaluation, we conduct aspect annotation of those reviews. To this end, we first (i) introduce an aspect typology and (ii) perform human and automatic annotation.

3.2.1 ASPECT TYPOLOGY AND POLARITY

We define a typology that contains 8 aspects, which follows the ACL review guidelines¹⁵ with small modifications, which are *Summary* (SUM), *Motivation/Impact* (MOT), *Originality* (ORI), *Soundness/Correctness* (SOU), *Substance* (SUB), *Replicability* (REP), *Meaningful Comparison* (CMP) and *Clarity* (CLA). The detailed elaborations of each aspect can be found in Supplemental Material B.1. Inside the parentheses are what we will refer to each aspect for brevity. To take into account whether the comments regarding each aspect are positive or negative, we also mark whether the comment is positive or negative for every aspect (except summary).

3.2.2 ASPECT ANNOTATION

Overall, the data annotation involves four steps that are shown in Fig. 1.

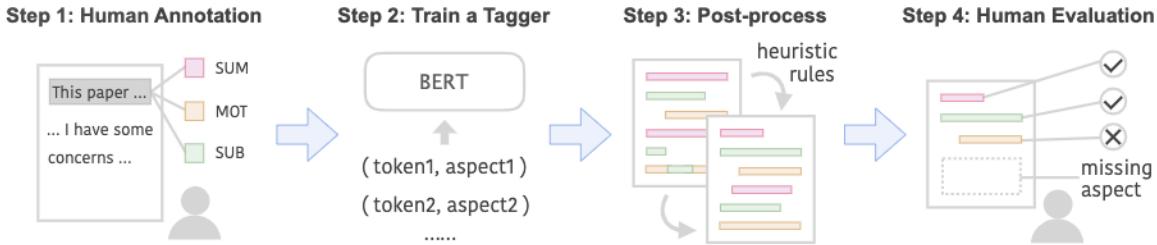


Figure 1: Data annotation pipeline.

Step 1: Manual Annotation To manually annotate aspects in reviews, we first set up a data annotation platform using Doccanno.¹⁶ We asked 6 students from ML/NLP backgrounds to annotate the dataset. We asked them to tag an appropriate text span that indicates a specific aspect. For example, “The results in this paper are new_[Positive Originality] and important to this field_[Positive Motivation]”. The detailed annotation guideline can be found in Supplemental Material B.1. Each review is annotated by two annotators and the lowest pair-wise Cohen kappa is 0.653, which stands for substantial agreement. In the end, we obtained 1,000 human-annotated reviews in total. The aspect statistics in this dataset are shown in Fig. 2-(a).

15. https://acl2018.org/downloads/acl_2018_review_form.html. We manually inspected several review guidelines from ML conferences and found the typology in ACL review guideline both general and comprehensive.

16. <https://github.com/doccano/doccano>

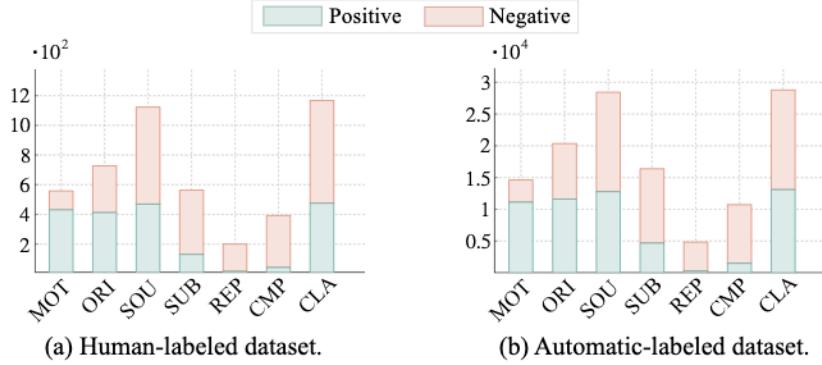


Figure 2: (a) and (b) represent distributions over seven aspects obtained by human and BERT-based tagger respectively. Green bins represent positive sentiment while red ones suggest negative sentiment. We omit “Sum” aspect since there is no polarity definition of it.

Step 2: Training an Aspect Tagger Since there are over 20,000 reviews in our dataset, using human labor to annotate them all is unrealistic. Therefore, we use the annotated data we do have to train an aspect tagger and use it to annotate the remaining reviews. The basic architecture of our aspect tagger contains a pre-trained model BERT (Devlin et al., 2019) and a multi-layer perceptron. The training details can be found in Appendix A.4.

Step 3: Post-processing However, after inspecting the automatically labeled dataset, we found that there appears to be some common problems such as interleaving different aspects and inappropriate boundaries. To address those problems, we used seven heuristic rules to refine the prediction results and they were executed sequentially. The detailed heuristics can be found in Appendix A.5. An example of our model prediction after applying heuristic rules is shown in Appendix A.6. Fig. 2-(b) shows the distribution of all reviews over different aspects. As can be seen, the relative number of different aspects and the ratio of positive to negative are very similar across human and automatic annotation.

Step 4: Human Evaluation To evaluate the data quality of reviews’ aspects, we conduct human evaluation. Specifically, we measure both aspect precision and aspect recall for our defined 15 aspects.

We randomly chose 300 samples from our automatically annotated dataset and assigned each sample to three different annotators to judge the annotation quality. As before, these annotators are all from ML/NLP backgrounds.

The detailed calculation for aspect precision and aspect recall can be found in Appendix A.7. Under these criteria, we achieved 92.75% aspect precision and 85.19% aspect recall. The fine-grained aspect precision and aspect recall for each aspect are shown in Tab. 3. The aspect recall for positive replicability is low. This is due to the fact that there are very few mentions of positive replicability. And in our human evaluation case, the system identified one out of two, which results in 50%. Other than that, the precision and recall are much higher.¹⁷

17. The recall numbers for negative aspects are lower than positive aspects. However, we argue that this will not affect the fidelity of our analysis much because (i) we observe that the imperfect recall is mostly

	SUM	MOT	ORI	SOU	SUB	REP	CLA	CMP
Polarity	+ -	+ -	+ -	+ -	+ -	+ -	+ -	+ -
Precision	95	94	72	95	94	95	90	100
Recall	100	89	71	87	80	98	79	94

Table 3: Fine-grained aspect precision and recall for each aspect. + denotes positive and – denotes negative.

Besides, one thing to mention is that our evaluation criterion is very strict, and it thus acts as a lower bound for these two metrics.

4. Scientific Review Generation

4.1 Task Formulation

The task of scientific review generation is similar in some ways to *scientific paper summarization*, which has been examined in some previous work, but review generation also has some important differences. Specifically, most current works summarize a paper (i) either from an “*author view*” that only use content written by the author to form a summary (Cohan et al., 2018a; Xiao & Carenini, 2019; Erera et al., 2019; Cohan et al., 2018a; Cachola et al., 2020b), (ii) or from a “*reader view*” that argues a paper’s summary should take into account the view of those in the research community (i.g., how other researchers think about this paper) (Cohan & Goharian, 2017; Yasunaga et al., 2019a).

In this work, we extend the view from “author” or “reader” to “reviewer.” When a reviewer comments upon a paper, they do not only reflect the core ideas in the original text but also add critical comments from different aspects, which usually requires external knowledge and domain expertise beyond the source paper itself. If we were able to automatically generate reviews similar to those created by competent reviewers, this could be useful to: (i) *authors*: helping them identify weak points in their paper and make it stronger. (ii) *reviewers*: improving the accuracy or comprehensiveness of the reviewing process by pointing out strong or weak points of the paper they may have missed. (iii) *readers*: helping them quickly grasp the main idea of the paper, and also pointing out parts of the paper that they should perhaps be skeptical of. The three views of scientific papers are shown in Fig. 3.

4.2 System Design

Despite the fact that our dataset contains fewer training samples compared with other benchmark summarization datasets, the few-shot learning ability of recent contextualized pre-trained models (Radford et al., 2019; Brown et al., 2020; Cachola et al., 2020a) still put training a passable review generation system from this dataset within grasp. We use BART (Lewis et al., 2019), which is a denoising autoencoder for pretraining sequence-to-sequence

(over 85%) caused by partial recognition of the same negative aspect in a review instead of inability to recognize at least one. This will not affect our calculation of *Aspect Coverage* and *Aspect Recall* very much. (ii) The imperfect recall will slightly pull up *Aspect Score* (will discuss in §5.3.1), but the trend will remain the same.

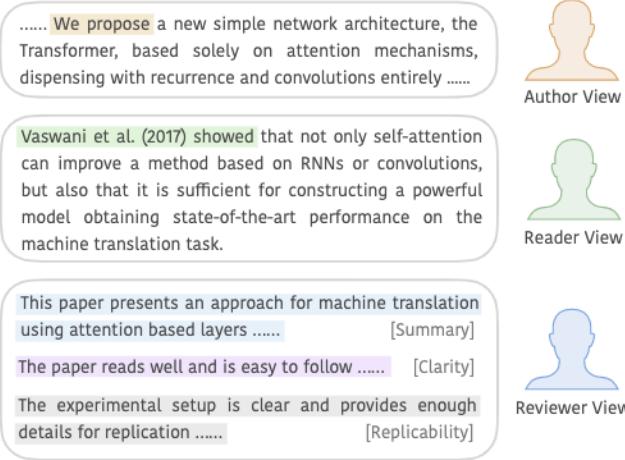


Figure 3: Three different views, “author,” “reader,” “reviewer,” for the paper “Attention Is All You Need” (Vaswani et al., 2017) from the paper’s abstract, citation (i.e., a paper that cites this paper) and peer review respectively.

models, as our pre-trained model since it has shown superior performance on multiple generation tasks.

However, even if we can take the advantage of this pre-trained model, how to deal with lengthy text in the context of using a pre-trained model (BART, for example, has a standard length limit of 1024 since it was pre-trained on texts of this size) remains challenging. After multiple trials, we opted for a two-stage method detailed below, and describe other explorations that were less effective in Appendix A.8.

4.2.1 TWO-STAGE SYSTEMS FOR LONG DOCUMENTS

Instead of regarding text generation as a holistic process, we decompose it into two steps, using an *extract-then-generate* paradigm (Chen & Bansal, 2018; Gehrmann et al., 2018; Subramanian et al., 2019; Dou et al., 2020). Specifically, we first perform content selection, extracting salient text pieces from source documents (papers), then generate summaries based on these extracted texts.

To search for an effective way to select content that is most useful for constructing a review generation system, we operationalize the first extraction step in several ways. One thing to notice is that the extraction methods we use here mainly focus on heuristics. We leave more complicated selection methods for future work.

Oracle Extraction First, for comparison purposes, we construct an oracle for each paper which is the extraction that achieves the highest average ROUGE score (i.e., the average score of ROUGE-1, ROUGE-2 and ROUGE-L) with respect to reference reviews, specifically using the greedy method described in Nallapati et al.’s (2017). Note that for each paper with multiple reviews, we construct multiple oracles for that paper. We assume that oracle extractions can reflect where reviewers pay more attention to when they are writing reviews. The selected sentence position distribution in oracles is shown in Fig. 4.

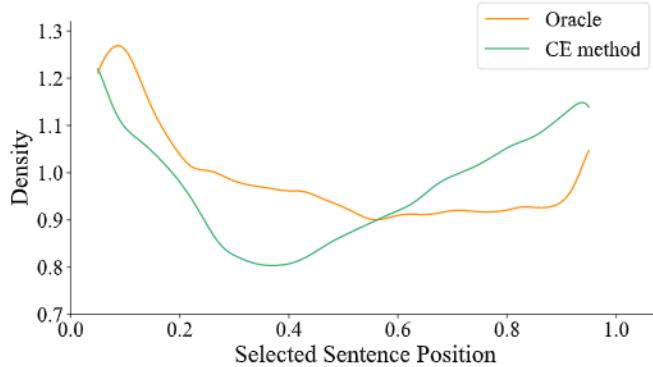


Figure 4: Selected sentence position distribution. We use the relative position of each sentence with regard to the whole article, thus taking values from 0 to 1.

Section-based Extraction Scientific papers are highly structured. As a convention, a scientific paper usually describes problem background, related work comparison, as well as its own contributions in the introduction part. Regarding this method, we only use the introduction section as our extraction for training, which can be regarded as a baseline model.

Cross-entropy (CE) Method Extraction Here we select salient sentences from the full text range. The way we do so is through a two-step selection process:

1. Select sentences containing certain informative keywords (e.g. propose) which are detailed in Appendix A.9. Those selected sentences form a set \mathcal{S} .
2. Select a subset $\mathcal{S}' \subseteq \mathcal{S}$ such that sentences in \mathcal{S}' cover diverse content and satisfy a length constraint.

In the second step, we use the cross-entropy method introduced in Feigenblat et al.’s (2017) where we select diverse content by maximizing unigram entropy. The details of this two-step process can be found in Appendix A.9. The selected sentence position distribution using this method is shown in Fig. 4. We can see that the extractor tends to select sentences from the beginning of a paper as well as the ending part of a paper just as the oracle extractor does. This makes sense because the beginning part is the introduction part which talks about the essence of the whole paper and the ending part mostly contains the analysis of experimental results and conclusions etc.

Hybrid Extraction We combine the abstract of a paper and its CE extraction to form a hybrid of both.

4.2.2 ASPECT-AWARE SUMMARIZATION

Typically in the *extract-then-generate* paradigm, we can just use the extractions directly and build a sequence-to-sequence model to generate text. Here, in order to generate reviews with more diverse aspects and to make it possible to interpret the generated reviews through the lens of their internal structure, we make a step towards a generation

framework involving *extract-then-generate-and-predict*. Specifically, unlike existing aspect-based summarization works that explicitly take aspects as input (Angelidis & Lapata, 2018; Frermann & Klementiev, 2019; Hayashi et al., 2020), we use our annotated aspects (§3.2) as additional information, and design an auxiliary task that aims to predict aspects of generated texts (reviews). Fig. 5 illustrates the general idea of this.

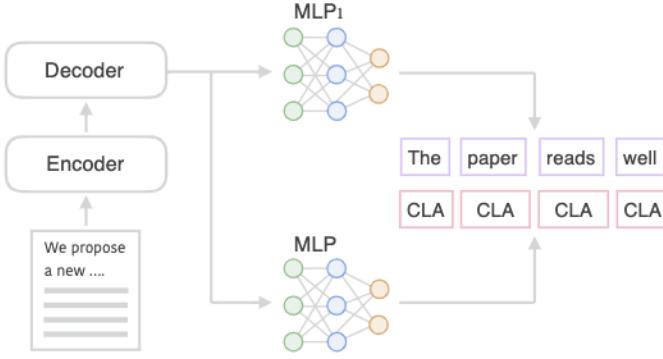


Figure 5: Aspect-aware summarization.

The loss of this model is shown in Eq. 3

$$\mathcal{L} = \mathcal{L}_{\text{seq2seq}} + \alpha \mathcal{L}_{\text{seqlab}} \quad (3)$$

where $\mathcal{L}_{\text{seq2seq}}$ denotes sequence to sequence loss which is the negative log likelihood of the correct next tokens, and $\mathcal{L}_{\text{seqlab}}$ denotes sequence labeling loss which is the negative log likelihood of the correct labels of next tokens. α is a hyper-parameter ($\alpha = 0.1$) that is tuned to maximize aspect coverage on the development set.

5. Experiment

In this section, we investigate using our proposed review generation systems with state-of-the-art pre-trained models, to what extent can we realize desiderata of reviews that we defined in §2.2. We approach this goal by two concrete questions: (1) *What are review generation systems (not) good at?* (2) *Will systems generate biased reviews?*

5.1 Settings

Here we consider three extraction strategies in §4.2.1 as well as two generation frameworks, one is the vanilla sequence to sequence model, the other is jointly sequence to sequence and sequence labeling.

Dataset We use our constructed dataset **ASAP-Review** described in §3 to conduct experiments. For each paper, we use full text (without Appendix) as source document.¹⁸ And we filtered papers with full text fewer than 100 words since they don't contain enough

¹⁸. If a paper has more than 250 sentences, we truncate it and take the first 250 sentences when we do the extraction step.

information for models to learn. For reviews, we only use 100-1024 word reviews¹⁹ for training due to computational efficiency, which account for 92.57% of all the reviews. This results in 8,742 unique papers and 25,986 paper-review pairs in total, the split of our dataset is shown in Tab. 4.

	Train	Validation	Test
Unique papers	6,993	874	875
Paper-review pairs	20,757	2,571	2,658

Table 4: Data split of **ASAP-Review**.

Model As mentioned in §4.2, the pre-trained sequence-to-sequence model we used is BART. For all models, we initialized the model weights using the checkpoint: “`bart-large-cnn`” which is pre-trained on “CNN/DM” dataset (Hermann et al., 2015).²⁰ For *extract-then-generate-and-predict* framework, we add another multilayer perceptron on top of the BART decoder, and initialize it with 0.0 mean and 0.02 standard deviation. We use the Adam optimizer(Kingma & Ba, 2014) with a linear learning rate scheduler which increases the learning rate linearly from 0 to $4e^{-5}$ in the first 10% steps (the warmup period) and then decreases the learning rate linearly to 0 throughout the rest of training steps. We finetuned our models on the whole dataset for 5 epochs. We set a checkpoint at the end of every epoch and finally took the one with the lowest validation loss.

During generation, we used beam search decoding with beam size 4. Similarly to training time, we set a minimum length of 100 and a maximum length of 1024. A length penalty of 2.0 and trigram blocking (Paulus et al., 2017) were used as well.

5.2 What are Systems Good and Bad At?

Based on the evaluation metrics we defined in §2.2, we conduct both automatic evaluation and human evaluation to characterize both reference reviews and generated reviews, aiming to analyze what sub-tasks of review generation automatic systems can do passably at, and also where they fail. The aspect information in each review is obtained using aspect tagger we trained in §3.2.

Automatic Evaluation Automatic evaluation metrics include *Aspect Coverage* (ACov), *Aspect Recall* (AREC) and *Semantic Equivalence* (ROUGE, BERTScore). Notably, for each source input, there are multiple reference reviews. When aggregating ROUGE and BERTScore²¹, we take the maximum instead of average. And when aggregating other metrics for human reviews, we take the average for each source document. The results are shown in Tab. 5.

Human Evaluation Metrics that require human labor include *Recommendation Accuracy* (RAcc), *Informativeness* (INFO), *Aspect-level Constructiveness* (ACON) and *Summary*

19. As measured by BART’s subword tokenizer.

20. We also tried “`bart-large-xsum`” checkpoint which is pre-trained on “XSUM dataset (Narayan, Cohen, & Lapata, 2018)”, however that results in much shorter reviews, and sentences in it tend to be succinct.

21. We have used our own custom baseline to rescale BERTScore, details can be found in Appendix A.3.

Accuracy (SAcc). We select 28 papers from ML/NLP/CV/RL domains. None of these papers are in the training set. Details regarding human judgment are shown in Appendix A.3. The evaluation results are shown in Tab. 5.

Desiderata	Decisive.	Comprehen.	Justification	Accuracy		Others					
Metric	RAcc	ACov	ARec	Info	ACon	SAcc	R-1	R-2	R-L	BS	
HUMAN	30.32	49.85	58.66	97.97	75.67	90.77	—	—	—	—	
EXTRACTIVE											
INTRO	—	—	—	—	—	—	38.62	8.84	25.11	29.22	
CE	—	—	—	—	—	—	38.56	7.81	25.94	29.11	
ABSCE	—	—	—	—	—	—	37.55	8.53	25.85	31.99	
EXTRACTIVE+ABSTRACTIVE											
Aspect											
INTRO	✗	-15.38†	50.37	55.52†	100.00†	43.78†	83.93	41.39	11.53	38.52	42.29
	✓	-11.54†	51.50	58.24	99.29	32.51†	80.36†	41.31	11.41	38.38	42.33
CE	✗	-23.08†	62.64†	60.73	99.29	39.17†	78.57†	42.37	11.72	39.86	41.78
	✓	-30.77†	63.96†	61.62†	100.00†	34.46†	69.64†	42.27	11.62	39.73	41.71
ABSCE	✗	-30.77†	55.37†	58.31	98.21	34.75†	92.86	43.11	12.24	40.18	42.90
	✓	-38.46†	56.91†	57.56	98.21	35.21†	87.50	42.99	12.19	40.12	42.63

Table 5: Results of the baseline models as well as different aspect-enhanced models under diverse evaluation metrics. “BS” represents BERTScore. † denotes that the difference between system generated reviews and human reviews are statistically significant (p-value < 0.05 using 10,000 paired bootstrap resampling (Efron, 1992) tests with 0.8 sample ratio).

Overall, we make the following observations:

5.2.1 WEAKNESSES

Review generation system will generate non-factual statements for many aspects of the paper assessment, which is a serious flaw in a high-stakes setting.

Lacking High-level Understanding Specifically, when using metrics that require higher level understanding of the source paper like *Recommendation Accuracy* and *Aspect-level Constructiveness*, proposed systems achieved much lower performance, with even the smallest gaps between systems and humans being 41.86% for *Recommendation Accuracy* and 31.89% for *Aspect-level Constructiveness* compared to reference reviews. This means our systems cannot precisely distinguish high-quality papers from low-quality papers and the evidence for negative aspects is not reliable most of the time.²²

Imitating Style After careful inspection, we find that some of sentences will appear frequently in different generated results. For example, “The paper is well-written and easy to follow” appears in more than 90% of generated reviews due to the fact that in

22. Although there exist varying degrees of performance differences on RAcc and ACon for different systems, we only find one pair of systems perform statistically different on ACon.

the training data, this exact sentence appears in more than 10% of papers. This suggests that the style of generated reviews tend to be influenced by high-frequency sentence patterns in training samples.

Lack of Questioning Generated reviews ask few questions about the paper content, which is an important component in peer reviewing. In the reference reviews, the average number of questions per review is 2.04, while it is only 0.32 in generated reviews.

5.2.2 ADVANTAGES

We find that review generation systems can often precisely summarize the core idea of the input paper, and generate reviews that cover more aspects of the paper’s quality than those created by human reviewers. Systems with aspect information are also aspect-aware and evidence-sensitive as we will discuss below.

Comprehensiveness In terms of *Aspect Coverage* and *Informativeness*, our systems can outperform human reviewers by at most 14.11% and 2.03% respectively, suggesting that even reviews from the reviewers may also fall short on our defined criterion regarding comprehensiveness.

Good Summarization Current systems can correctly summarize the contributions of papers most of the time as shown by *Summary Accuracy*. 4 out of 6 systems can achieve over 80% accuracy and statistical significance tests show that gaps between top-3 systems and human reviewers are not significant. This means that in terms of summarizing the paper content, current systems can achieve comparable performance to human reviewers.

5.2.3 SYSTEM COMPARISONS

We also look into how systems with different settings are diverse in performance and make the following observations.

Summarization Paradigms By looking at ROUGE (R) and BERTScore (BS), we can see that “extractive + abstractive”-based methods can consistently outperform pure extractive methods, with the smallest gaps of 2.69, 2.57, 12.44, 9.72 for R-1, R-1, R-L and BS respectively. This demonstrates the necessity of using abstractive summarization which can generate reviews that are close both in meaning as well as language use to human reviews.

Extraction Strategies We can see that it is more effective to use extracted text from the full paper to aid the generation process, resulting in higher aspect coverage compared with solely using introduction information. This is reasonable since models can obtain more diverse input from the full text.

System Diagnosis Our fine-grained evaluation metrics enable us to compare different systems and interpret their relative merits. For example, as discussed before, our systems can achieve higher *Informativeness* than reference reviews while suffering from much lower *Aspect-level Constructiveness*. This means that if we want our systems to match the performance of real reviewers, we should focus on improving the factuality of our systems instead of aiming for methods that provide more evidence for negative aspects (which are not factually correct most of the time in the first place).

5.2.4 CASE STUDY

To get an intuitive understanding of how aspect-enhanced review generation system worked, we perform analysis on a real case. (More analysis can be found in Appendix A.10.) Specifically, since our aspect-enhanced model is trained based on multi-task learning framework, it would be interesting to see how well the tokens are generated associated with corresponding aspects. We take our aspect-enhanced model using CE extraction to conduct this experiment. Tab. 6 shows an example review when we do so.

This paper studies the problem of transfer learning from a single pre-trained network onto a new task .
The authors propose a method of combining the outputs of multiple pre-trainable classifiers by training on their raw predictions and then fine-tuning on the target samples . The proposed method is based on the use of (non-linear) maximal correlation analysis that originated with Hirschfeld [9] to this problem . The paper is well_[Summary] written and easy to follow . The_[Positive Clarity] experimental results show that_[Positive Substance] the proposed method outperforms the state-of-the-art methods on the CIFAR-100, Stanford Dogs , and Tiny ImageNet datasets . However , it is not clear to me how the performance of the method is affected by_[Negative Soundness] the number of target training samples . It would be better if the authors_[Negative Substance] can provide some theoretical_[Negative Substance] analysis on the_[Negative Substance] effect of the size of the_[Negative Substance] target dataset .

Table 6: Illustration of generated tokens associated with corresponding aspects. + denotes positive sentiment. - denotes negative sentiment.

We can see that the model can not only generate fluent text but also be aware of what aspect it is going to generate as well as the correct polarity of that aspect. Although the generated aspects are often small segments and there are some minor alignment issues, the model is clearly aspect-aware.

5.3 Will System Generate Biased Reviews?

Biases in text are prevalent, but often challenging to detect (Manzoor & Shah, 2020; Stelmakh et al., 2019). For example, in natural language processing, researchers are trying to identify societal biases (e.g, gender) in data and learning systems on different tasks (Bolukbasi et al., 2016; Zhao et al., 2018; Stanovsky et al., 2019). However, previous works on analyzing bias in scientific peer review usually focus on disparities in numerical feedback instead of text. Manzoor and Shah’s (2020) recently uncover latent bias in peer review based on review text. In this work, besides designing a model to generate reviews, we also perform an analysis of disparity, in which we propose a method to identify and quantify disparities both in human-labeled and system-generated data in a more fine-grained fashion. One thing to notice is that disparities in reviews do not necessarily indicate the presence of bias, as it is difficult to control for the quality of papers in different groups.

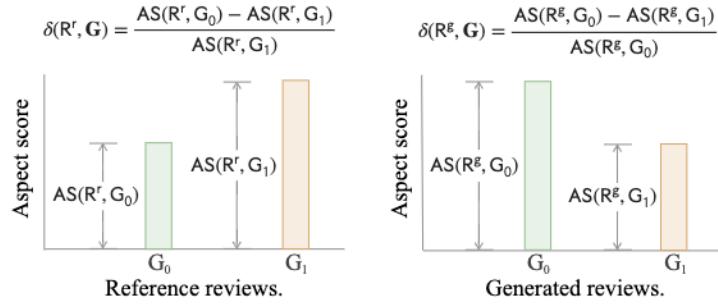


Figure 6: *Aspect score* $AS(R, G_i)$ and *disparity* $\delta(R, \mathbf{G})$ in reference reviews (R^r) and generated reviews (R^g). $\mathbf{G} = [G_0, G_1]$ denotes different groups.

5.3.1 MEASURING DISPARITY IN REVIEWS

To characterize disparities existing in reviews, we (i) first define an *aspect score*, which calculates the percentage of positive occurrences²³ of each aspect. The polarity of each aspect is obtained based on our learned tagger in §3.2; (ii) then we aim to observe if different groups G_i (e.g., groups whether the paper is anonymous during reviewing or is not anonymous) of reviews R would exhibit *disparity* $\delta(R, \mathbf{G})$ in different aspects. The calculation of disparity can be visualized in Fig. 6.

Based on above two definitions, we characterize disparity in two ways respectively:

- (1) **spider chart**, which directly visualizes aspect scores of different groups of reviews w.r.t each aspect.
- (2) **disparity difference**, which represents the difference between disparities in generated reviews R^g and reference reviews R^r and can be formally calculated as:

$$\Delta(R^g, R^r, \mathbf{G}) = \delta(R^g, \mathbf{G}) - \delta(R^r, \mathbf{G}) \quad (4)$$

where $\mathbf{G} = [G_0, G_1]$ denotes different groups based on a given partition criterion. Positive value means generated reviews favor group G_0 more compared to reference reviews, and vice versa.

In this work, we group reviews from two perspectives. The basic statistics are shown in Tab. 7.

	Native	Non-native	Anonym.	Non-anonym.
Total	651	224	613	217
Acc.%	66.51%	50.00%	57.59%	78.34%

Table 7: Test set statistics based on nativeness and anonymity.

23. If an aspect does not appear in a review, then we count the score for that aspect 0.5 (stands for neutral). Details see Appendix A.11.

Nativeness We categorize all papers in test set into “native” (G_0) and “non-native” (G_1) based on whether there is at least one native speaker in the author list as well as whether the institution is in an English-speaking country.²⁴

Anonymity We categorize all papers in test set into “anonymous” (G_0) and “non-anonymous” (G_1) based on whether the paper has been released as a pre-print before a half month after the conference submission deadline.²⁵

Here we take our model with introduction extraction as an example to showcase how to use the fine-grained aspect information in our dataset to do disparity analysis. We list the disparity analysis for other models in Appendix A.12.

5.3.2 NATIVENESS ANALYSIS

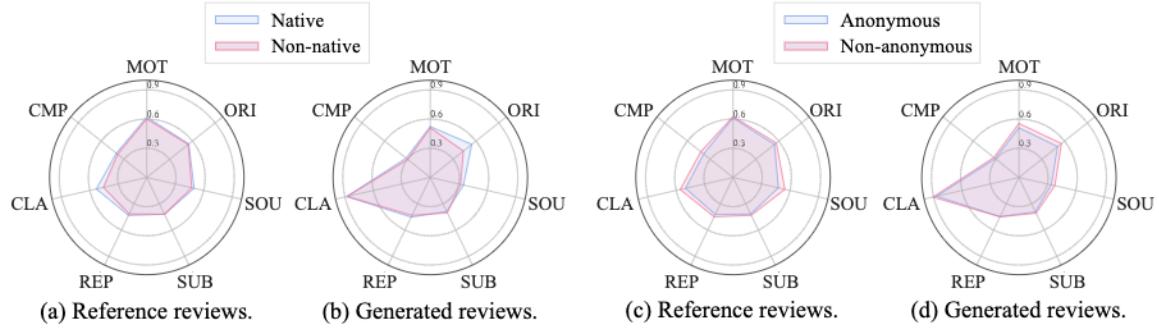


Figure 7: Spider chart of aspect scores with respect to different groups.

Spider Chart Generally, Native papers receive higher score in most aspects in both reference reviews and generated reviews. Specifically, for human reviews: (1) By looking at Fig. 7-(a), there is a significant gap in Clarity, which is reasonable since non-native authors may have more trouble conveying their ideas. (2) Scores of the two groups are much closer in other aspects. For system-generated reviews: As shown in Fig. 7-(b), the auto-review system narrows the disparity in Clarity but amplifies it in Originality, meaning that system reviewers are harsher than human reviewers when commenting the paper’s “Originality” for non-native English speakers. This observation suggests that a review system **might** generate biased reviews in some aspects, which could lead to unfair comments. Therefore, a system should be de-biased before it comes to use.

Disparity Difference Through spider chart, gaps between different groups are relatively small and hard to discern. Besides, those gaps can only show the absolute favor for a certain group in different aspects. We are also interested in whether generated reviews are more in favor of a certain group **compared to reference reviews**. To do this, we calculate disparity differences and list them in Tab. 8.

24. We used <https://www.familysearch.org/en/> to decide the nationality of an author. In cases where all authors are not from an english-speaking country, we look into the institution information to further decide the attribution of the paper based on whether the institution is from an english-speaking country.

25. We discard papers from ICLR 2017 since the reviewing process was single blind.

	MOT	ORI	SOU	SUB	REP	CLA	CMP	Total
Nativeness	-0.72	+18.71	+3.84	-3.66	+0.73	-13.32	+2.40	43.39
Anonymity	-5.69	-4.43	+2.76	-0.64	+5.65	+5.80	+3.02	28.00

Table 8: Disparity differences regarding nativeness and anonymity. Total is the sum of absolute value of disparity difference.

As shown in Tab. 8, for **Originality** and **Clarity**, the disparity difference is +18.71 and -13.32 which means that the system favours native papers in **Originality** and non-native papers in **Clarity compared to** human reviewers. This observation is consistent with spider chart.

5.3.3 ANONYMITY ANALYSIS

Spider Chart By looking at Fig. 7-(c) and Fig. 7-(d), we find that both human reviewers and system reviewers favor non-anonymous papers in all aspects. Specifically, for human reviews: we find gaps are non-negligible in **Soundness**, **Clarity** and **Meaningful Comparison** while for system-generated reviews, we observe that gaps are considerable in **Motivation**, **Originality**, **Soundness**. This observation is interesting since human reviewers may be aware of the identity of the authors due to non-anonymity which may affect the reviews they write. However, our system is not aware of that and its preference towards non-anonymous paper probably suggests some quality difference.²⁶

Disparity Difference By looking at Tab. 8, we find that the largest absolute disparity difference regarding anonymity is 5.80 compared to 18.71 regarding nativeness. This suggests that regarding anonymity, our system’s preference does not diverge that much from human reviewers. Also, the total aspect bias regarding anonymity is 28.00, much smaller compared to total aspect bias regarding nativeness (43.00). This also suggests that our model is less sensitive to anonymity compared to nativeness.

The observations above are probably related to some superficial heuristics existing in peer review. For example, when reviewers detect some grammar mistakes, they may assume that the authors are not native and then bias towards rejecting the paper by claiming some clarity issues. Another example is that there may exist differences in the research topics pursued by different subgroups (e.g., different countries), the bias regarding nativeness may also suggest the favor of certain topics in the reviewing process. Those superficial heuristics should be discouraged and deserve further investigation in future research.

6. Related Work

Scientific Review Generation There has been a relative paucity of work on scientific review generation, other than Bartoli et al.’s (2016) investigating the feasibility of generating fake reviews by surface-level term replacement and sentence reordering etc. In addition contemporaneous and independent work by Wang et al.’s (2020a) proposes a two-stage

26. Non-anonymous papers are more likely to have been rejected before and therefore are revised many more times.

information extraction and summarization pipeline to generate paper reviews. Their evaluation focuses mainly on the accuracy of information extraction, and the evaluation of the generated reviews is somewhat precursory, assessing only a single criterion “constructiveness and validity” manually over 50 papers. Our paper (1) proposes a wide variety of diagnostic criteria on review quality, (2) uses a very different summarization methodology, and (3) evaluates the generated results extensively.

Peer Review Peer review is an essential component of the research cycle and is adopted by most journals and conferences to identify important and relevant research. However, at the same time it is easy to identify many issues: expensiveness, slowness, existence of inconsistency (Langford & Guzdial, 2015) and bias (Tomkins et al., 2017), etc.

Some efforts have been put into analyzing the peer review process including automating review assignment (Jin et al., 2017; Nguyen et al., 2018; Anjum et al., 2019; Jecmen et al., 2020), examining bias problems (Tomkins et al., 2017; Stelmakh et al., 2019), examining consistency problems (Langford & Guzdial, 2015) and performing sentiment analysis on reviews (Wang & Wan, 2018; Chakraborty et al., 2020). Several decision classification methods have been explored to help make accept or reject decision given a paper. Those methods are either based on textual (Kang et al., 2018; Qiao et al., 2018) or visual (Von Bearnensquash, 2010; Huang, 2018) information. However, they do not directly alleviate review load, as our paper aims to do.

7. Discussion and Future Directions

We first summarize what we have achieved in this work and discuss how the current *ReviewAdvisor* system may potentially help in a paper authoring or reviewing process. Then we discuss challenges and potential directions for the automatic review generation task, which, hopefully, encourages more future researchers to explore this task, and in the right direction.

7.1 Machine-assisted Reviewing and Authoring

Reviewing Advisor Given the current quality of the generated reviews, it is most likely better to consider whether *ReviewAdvisor* may have utility in a system to assist human authors or reviewers in the machine learning domain. Although there is still a large room for improvement, our results indicate that even with current technology: (1) Based on the evaluation of §5.2, *Summary Accuracy* of our systems is quite high, suggesting that it could be used general readers or reviewers alike to quickly understand the core ideas of submitted or published papers. (2) Based on evaluation of §5.2, reviews generated by *ReviewAdvisor* can cover more aspects and generate more informative reviews. Although the associated opinions often are not factual, they may still provide a template of what content to include in a review or what phrasing to use when describing particular aspects. In order to prevent reviewers from quickly agreeing to generated reviews, we have also taken some measures. (1) For the same paper, we provide multiple reviews generated by different systems for reference. (2) For each aspect (e.g., *Clarity*), our systems can provide relevant evidence sentences from the paper, helping reviewers quickly identify salient information when reviewing the paper (Detailed example in our Appendix A.10). (3) We also remind users on the usage page of

our systems that the generated reviews are for reference only, and users should have their own judgment.

Revision Advisor Another potential role of such a system is to use it as an assistance system for authoring. These generated critical reviews may help researchers identify the shortcomings (e.g., grammatical errors, missing references) of their preliminary drafts and guide them on how to do paper revision. Of course factuality is an issue here as well, but even with the current technology, authors can decide which comments to pay attention to and which to ignore.

7.2 Challenges and Promising Directions

7.2.1 MODEL

(1) *Long Document Modeling*: The average length of one scientific paper is commonly larger than 5,000 words, far beyond the input text’s length that mainstream neural sequence models (e.g., LSTM, Transformer) or pre-trained models (e.g., BERT, BART) normally use. This work (in §4.2.1) bypasses the difficulty by using a two-stage system, but other strategies should be explored.

(2) *Pre-trained Models for the Scientific Domain*: Although previous works, as exemplified by (Beltagy et al., 2019) have pre-trained BERT on scientific domain, we observe that using these models with transformer decoders perform much worse than BART on sequence generation tasks in terms of fluency and coherence, which calls for general sequence-to-sequence models pre-trained on scientific domain for higher-quality review generation.

(3) *Structure Information*: Review generation systems could get a deeper understanding of a given research paper if structural information can be provided. To this end, outputs from scientific paper-based information extraction tasks (Hou et al., 2019; Jain et al., 2020) could be utilized to guide review generation.

(4) *External Knowledge*: Besides the paper itself, human reviewers also rely on external knowledge of the research field as a whole. The lack of consideration of the surrounding literature is an obvious shortcoming of the current prototype system considered here. There are a number of different ways that the surrounding literature could be considered in an automatic review system, such as retrieval-based models (Guu et al., 2020; Lewis et al., 2020) that encode similar papers, citation graphs constructed based on encoding the citation relationships with other papers (Luu et al., 2020; Wang et al., 2020b; Lo et al., 2020), or knowledge graphs connecting concepts across different papers (Luan et al., 2018; Koncel-Kedziorski et al., 2019). Also, recently, August et al.’s (2020) compile a set of writing strategies drawn from a wide range of prescriptive sources, it would be also valuable to transfer this knowledge into the auto-review system.

7.2.2 DATASETS

(5) *More Open, Fine-grained Review Data*: In this work, we annotate fine-grained information (aspects) of each review manually. However, this information could potentially be obtained directly from the peer review system. How to access this information appropriately would be an important and valuable step in the future.

(6) *More Accurate and Powerful Scientific Paper Parsers:* Existing parsing tools (e.g. science-parse, grobid) for scientific papers are commonly designed for certain specific paper templates, and also still struggle at extracting fine-grained information, such as the content of tables and figures.

7.2.3 EVALUATION

(7) *Fairness and Bias in Generated Text:* In this work, we make a step towards identifying and quantifying disparities existing in human and system-generated reviews. Future works can explore more along this direction based on our dataset that contains fine-grained aspect annotation.

(8) *Factuality and Reliability:* A generated review should be *factually correct* (Wadden et al., 2020) which also poses challenge to the current evaluation methodology. In addition to generating a review, a reliable system should also provide a level of confidence with respect to the current comment. Moreover, whether review scores are calibrated is another valuable question.

7.3 Conclusion

In answer to the titular question of “can we automate scientific review,” the answer is clearly “not yet”. However, we believe the models, data, and analysis tools presented in this paper will be useful as a starting point for systems that can work in concert with human reviewers to make their job easier and more effective.

Acknowledgments

This work could not be accomplished without the help of many researchers. We would like to thank people for their generous support, especially,

We would like to thank the many researchers who volunteered to help us with human annotation and evaluation: Gábor Berend, Zhouhan Lin, William W. Cohen, Pengcheng Yin, Tiange Luo, Yuki M. Asano, Junjie Yan, Tuomas Haarnoja, Dandan Guo, Jie Fu, Lei Chen, Jinlan Fu, Yiran Chen, Jiapeng Wu, Wenshan Wang, Ziyi Dou, Yixin Liu, Junxian He, Bahetiyaer Bare, Saizheng Zhang, Jiateng Xie, Spyros Gidaris, Marco Federici, Junqi Dai, Zihuiwen Ye, Jie Zhou, Yufang Liu, Yue Zhang, Ruifeng Xu, Zhenghua Li, Chunting Zhou, Yang Wei, Jiahao Wang, Bowen Tan, Anda Zhou.

We would also like to thank Emily Bender, Peyman Milanfar, and Nihar Shah for their constructive comments on earlier versions of the paper draft. Some of the comments would certainly be given *negative* polarity according to our typology in Section 3.2.1, but they were very helpful nonetheless.

Appendix A. Details of Experiments

A.1 Review of this Paper Written by Our Model

Notably, the following review is generated based on the arXiv version of this paper without the TL;QR section (The original version can be found here: <https://drive.google.com/file/d/1nC4kCaaeqKRiajcvK75g421Ku9Jog1n9/view?usp=sharing>). And we directly put the system output here without any grammar check.

Summary : This paper presents an approach to evaluate the quality of reviews generated by an automatic summarization system for scientific papers . The authors build a dataset of reviews , named ASAP-Review1 , from machine learning domain , and make fine-grained annotations of aspect information for each review , which provides the possibility for a richer evaluation of generated reviews . They train a summarization model to generate reviews from scientific papers , and evaluate the output according to our evaluation metrics described above . They propose different architectural designs for this model , which they dub ReviewAdvisor , and comprehensively evaluate them , interpreting their relative advantages and disadvantages . They find that both human and automatic reviewers exhibit varying degrees of bias regarding native English speakers vs non-native English speakers , and find that native speakers tend to obtain higher scores on “ Clarity ” and “ Potential Impact ” . The paper is well-written and easy to follow .

Strengths : 1 . The proposed approach is novel and interesting . 2 .The paper is easy to read and well-organized . 3 .The evaluation metrics are well-motivated . 4 .The authors have done a good job of evaluating the proposed approach .

Weaknesses : 1) The evaluation metrics used in this paper are not well-defined . For example , what is the definition of “ good review quality ” ? What is the criteria for a good review ? 2) It is not clear to me how the authors define “ factually correct ” , “ fair ” or “ non-factual ” in Section 3.2 . 3) The authors should provide more details about the evaluation metrics in the paper . For instance , what are the criteria used in Table 1 and Table 2 ? What are the metrics used for the evaluation in Table 3 and Table 4 ? 4) It would be better if the authors can provide more explanations about the results of Table 2 and Table 3 . 5) In Table 3 , the authors mentioned that “ we found the constructed automatic review system generates nonfactual statements regarding many aspects of the paper assess- 1ASpect-enhanced-Anced Peer Review dataset , which is a serious flaw in a high-stakes setting such as reviewing . However , there are some bright points as well . ” However , it is unclear to me why the authors found this problem . 6) In Section 4.3 , it seems that the authors did not provide any explanation about why the human reviewers are biased . 7) In Figure 1 , it would be good to provide more information about the training data . 8) In section 4.4 , it will be better to provide some explanation about how the human reviews are generated .

A.2 Screenshot of Our Demo System

The screenshot of our Demo system is shown in Fig. 8.

Paper Review

[Model List](#)

[Introduction](#)

[CE extraction](#)

[Abstract + CE](#)

[Abstract + CE \(with aspect\)](#)

summary clarity **motivation** originality soundness substance

This paper presents a method for training deep neural networks with residual layers. The idea is to learn a residual representation for each layer of the network, which is learned as an identity mapping of the input and output channels. The paper is well written and easy to follow. The method is evaluated on ImageNet and CIFAR-10. The results show that the proposed method is able to improve the performance of deep networks with a depth of up to 152 layers (8x deeper than VGG nets). The authors also show that their method can achieve a 28 % improvement on the COCO object detection dataset. I think this paper is a good contribution to the field of deep learning. However, I have the following concerns: 1. The novelty of the method is limited. It is a simple idea to learn the residual representation of each layer as a mapping from the input to the output channels, and it is not very novel. 2. The experiments are not very convincing. The proposed method does not outperform VGG on the ImageNet dataset, and the results on Cifar-10 are not impressive. 3. The authors claim that the method can be applied to other vision and non-vision problems, but the experiments are only conducted on the image classification task. It would be better if the authors can conduct experiments on other vision tasks.

Paper Content

[Abstract](#)

[1. Introduction](#)

[2. Related Work](#)

[3.1. Residual Learning](#)

[3.2. Identity Mapping by
Shortcuts](#)

[3.3. Network Architectures](#)

[3.4. Implementation](#)

Deeper neural networks are more difficult to train. We present a residual learning framework to ease the training of networks that are substantially deeper than those used previously. We explicitly reformulate the layers as learning residual functions with reference to the layer inputs, instead of learning unreference functions. We provide comprehensive empirical evidence showing that these residual networks are easier to optimize, and can gain accuracy from considerably increased depth. On the ImageNet dataset we evaluate residual nets with a depth of up to 152 layers—8x deeper than VGG nets [40] but still having lower complexity. An ensemble of these residual nets achieves 3.57% error on the ImageNet test set. This result won the 1st place on the ILSVRC 2015 classification task. We also present analysis on CIFAR-10 with 100 and 1000 layers. The depth of representations is of central importance for many visual recognition tasks. Solely due to our extremely deep representations, we obtain a 28% relative improvement on the COCO object detection dataset. Deep residual nets are foundations of our submissions to ILSVRC & COCO 2015 competitions, where we also won the 1st places on the tasks of ImageNet detection, ImageNet localization, COCO detection, and COCO segmentation.

Figure 8: Generated reviews with aspects by our *ReviewAdvisor* Demo associated with evidence sentences from the paper “Deep Residual Learning for Image Recognition” (He et al., 2016).

A.3 Details for Evaluation Metrics

Rec In §2.2, the REC function we define corresponds to the recommendation sentiment of a review, with $\{-1, 0, 1\}$ representing negative, neutral and positive.

To decide the sentiment of a reference review, we use the rating information from reviewers: (i) if the rating corresponds to marginal accept or marginal reject, then we regard it as neutral; (ii) if the rating is above marginal accept, then we regard it as positive; (iii) otherwise, we regard it as negative.

To decide the sentiment of a generated review, two members of the project team judged the sentiment polarity of a review. If they agreed with each other, then we use the agreed-upon sentiment, if they disagreed with each other, then we label the sentiment of that review as neutral. The Cohen kappa of two annotators is 0.5778, which is commonly referred to as “moderate” agreement.

Info The judgement of evidence for negative aspects was conducted by a member of the project team, who judged whether each identified negative aspect was accompanied by evidence irrespective of the correctness of the evidence. In other words, as long as there is a reason, we count that as an evidence.

SAcc & ACon The judgement of summary accuracy and valid support for negative aspects are performed by one of the first authors of the reviewed paper. Summary and each negative aspect in the review should be scored 0, 0.5 or 1 which represent agree, partially agree and disagree respectively. We provide authors with the following instructions:

“We have created a Google doc for your paper, and you can score the summary as well as each aspect with its corresponding comments inside the red brackets. “1” denotes agree, “0.5” denotes partially agree, “0” denotes disagree. You only need to assign a score based on your judgment. For summary, agree means that you think it’s factually correct. For each aspect, agree means that you think the strength/weakness the reviewer points out is reasonable or constructive.”

BERTScore Regarding BERTScore, we apply the same rescaling procedure following (Zhang et al., 2019), which does not affect the ranking ability of BERTScore, but make the scores more discriminative.

A.4 Training of Aspect Tagger

We formulate the annotation process as a sequence labeling problem where the input is a sequence consisting of n words $S = w_1, \dots, w_n$, and the target is a sequence of tags one for each word $T = t_1, \dots, t_n$. We aim to find a mapping f such that $T = f(S)$ can convey reasonable aspect information in the input sequence.

We first segment each review into multiple sentences and consider each sentence as an individual training example.²⁷ For a tokenized sequence contains n tokens (w_1, w_2, \dots, w_n) , we use BERT to get a contextualized representation for each token (e_1, e_2, \dots, e_n) , where

27. We also tried using larger context such as paragraph, but found out the results less satisfying since the model identified fewer aspects.

e_i represents the vector for i -th token. Then those contextualized representations can be used as features for token classification:

$$\mathbf{p}_i = \text{softmax}(\mathbf{W}e_i + \mathbf{b})$$

where \mathbf{W} and \mathbf{b} are tunable parameters of the multilayer perceptron. \mathbf{p}_i is a vector that represents the probability of token i being assigned to different aspects.

We use the negative log likelihood of the correct labels as training loss:

$$\mathcal{L} = - \sum_{t \in \mathcal{T}} \log \mathbf{p}_{tj}$$

where j is the label of token t , and \mathcal{T} denotes all the tokens.

We used 900 annotated reviews for training and 100 for validation which is equivalent to using 16,543 training data and 1,700 validation data since we consider sentence as the basic individual training sample. The initial BERT checkpoint we used is “bert-large-cased” which is the large version of BERT with a cased vocabulary. We used Adam optimizer (Kingma & Ba, 2014) with a learning rate of $5e^{-5}$ to finetune our model. We trained for 5 epochs and saved the model that achieved lowest loss on validation set as our aspect tagger.

A.5 Heuristics for Refining Prediction Results

The seven heuristic rules used for refining the prediction results are listed below. Examples of those rules are shown in Tab. 9.

1. If there are no other tags (they are tagged as “0” which stands for *Outside*) between two “summary” tags, then replace all tags between them with “summary” tag.
2. If there are multiple discontinuous text spans tagged as “summary”, we keep the first one and discard others.
3. If the punctuation is separately tagged and is different from its neighbors, we replace its tag to “0”.
4. If two identical tags are separated by a single other tag, then replace this tag with its right neighbor’s tag.
5. If there exists a single token with a tag and its neighbors are “0”, then replace this tag to ‘0’.
6. For a “non-summary” “non-0” tag span, if its neighbors are “0” and the start/end of this span is not special symbol (for example, punctuations or other symbols that have 1 length), then we extend from its start/end until we meet other “non-0” tag or special symbol.
7. If the “summary” span does not end with a period, then we truncate or extend it at most five words to make it ends with a period.

Heur.	Before	After
1	The authors present a method for learning Hamiltonian functions [summary] ... [o] this is [Summary] ... [o] that maps past observations to a latent p,q space in a VAE-like fashion. [Summary]	The authors present a method for learning Hamiltonian functions ... this is ... that maps past observations to a latent p, q space in a VAE-like fashion. [summary]
2	This paper proposes a new representation learning model for graph optimization, Graph2Seq .[Summary] ... the theorems are very interesting .[Positive Originality] ... The performance of Graph2Seq is remarkable. [Summary]	This paper proposes a new representation learning model for graph optimization, Graph2Seq .[Summary] ... the theorems are very interesting .[Positive Originality] ... The performance of Graph2Seq is remarkable. [o]
3	The proposed idea is novel [Positive Originality] .[Positive Motivation] The paper is well written and easy to follow.[Positive Clarity]	The proposed idea is novel [Positive Originality] .[o] The paper is well written and easy to follow.[Positive Clarity]
4	The overall notion of [Positive Originality] learning [o] a Hamiltonian network directly is a great one.[Positive Originality]	The overall notion of learning a Hamiltonian network directly is a great one. [Positive Originality]
5	It is [o] clearly [Positive Clarity] geared towards DNN practitioners. [o]	It is clearly geared towards DNN practitioners. [o]
6	In contrast , this aspect [o] is missing from other work on ML [Negative Meaningful Comparison] for optimization. [o]	In contrast, this aspect is missing from other work on ML for optimization. [Negative Meaningful Comp.]
7	The authors propose a novel approach to estimate unbalanced optimal transport between sampled measures that scales well in the dimension and in the number of samples ... The effectiveness of the approach [Summary] is shown on some tasks. [o]	The authors propose a novel approach to estimate unbalanced optimal transport between sampled measures that scales well in the dimension and in the in the number of samples ... The effectiveness of the approach is shown on some tasks. [Summary]

Table 9: Examples of seven heuristic rules used for refining prediction results.

A.6 An Example of Automatically Annotated Reviews

Tab. 10 illustrates an annotated review after using our trained aspect tagger and heuristic refining rules in Appendix A.5. Although here we do not add separate polarity tags to avoid visual burden, the polarity of each aspect the model predicts is correct.

A.7 Calculation of Aspect Precision and Aspect Recall

To measure aspect precision, we asked three annotators to decide whether each aspect span the model predicted is accurate. They were asked to delete a tagged span if they regarded

This paper studies the graph embedding problem by using encoder-decoder method . The experimental study on real network data sets show the features extracted by the proposed model is good for classification ._[Summary] Strong points of this paper: 1. The idea of using the methods from natural language processing to graph mining is quite interesting ._[Originality] 2. The organization of the paper is clear Weak points of this paper_[Clarity]: 1. Comparisons with state-of-art-methods (Graph Kernels) is missing._[Comparison] 2. The problem is not well motivated_[Motivation], are there any application of this. What is the difference from the graph kernel methods ? The comparison with graph kernel is missing ._[Comparison] 3. Need more experiment to demonstrate the power of their feature extraction methods ._[Substance] (Clustering, Search, Prediction etc.) 4. Presentation of the paper is weak . There are lots of typos and unclear statements. _[Clarity]

Table 10: An example of automatically labeled reviews.

it as inappropriate. We denote all prediction spans as \mathcal{M} , and the filtered spans from annotators as \mathcal{F}_1 , \mathcal{F}_2 and \mathcal{F}_3 . We represent $n_{\mathcal{S}}$ as the total number of text spans in \mathcal{S} . Here we define correct spans as

$$\mathcal{C} = \{l | l \in \mathcal{F}_1, l \in \mathcal{F}_2, l \in \mathcal{F}_3\}$$

The aspect precision is calculated using Formula 5.

$$\text{Precision} = \frac{n_{\mathcal{C}}}{n_{\mathcal{M}}} \quad (5)$$

For measuring aspect recall, we asked three annotators to label aspect spans that they identified while the model ignored. We denote the additional labeled spans from one annotator as \mathcal{A} where $\mathcal{A} = \{a_1, a_2, \dots, a_{n_{\mathcal{A}}}\}$, a_i represents a text span. We denote the additional labeled spans from other two annotators as \mathcal{B} and \mathcal{C} .

We define common ignored spans for every two annotators as below. $|\cdot|$ denotes the number of tokens in a span and \cap takes the intersect span between two spans.

$$\begin{aligned} \mathcal{I}_1 &= \{a_i \cap b_j | \frac{|a_i \cap b_j|}{\min \{|a_i|, |b_j|\}} > 0.5\} \\ \mathcal{I}_2 &= \{b_i \cap c_j | \frac{|b_i \cap c_j|}{\min \{|b_i|, |c_j|\}} > 0.5\} \\ \mathcal{I}_3 &= \{a_i \cap c_j | \frac{|a_i \cap c_j|}{\min \{|a_i|, |c_j|\}} > 0.5\} \end{aligned}$$

We also define common ignored spans for three annotators as below.

$$\mathcal{I} = \{a_i \cap b_j \cap c_k | \frac{|a_i \cap b_j \cap c_k|}{\min \{|a_i|, |b_j|, |c_k|\}} > 0.3\}$$

where a_i, b_j, c_k are text spans from $\mathcal{A}, \mathcal{B}, \mathcal{C}$ respectively. We assume all the spans the model predicts are correct. Then we can calculate total number of spans using Formula 6.

$$n = n_{\mathcal{M}} + n_{\mathcal{A}} + n_{\mathcal{B}} + n_{\mathcal{C}} - n_{\mathcal{I}_1} - n_{\mathcal{I}_2} - n_{\mathcal{I}_3} + n_{\mathcal{I}} \quad (6)$$

The aspect recall is calculated using Formula 7.

$$\text{Recall} = \frac{n_{\mathcal{M}}}{n} \quad (7)$$

A.8 Adjusting BART for Long Documents

KEYWORDS							
find	prove	examine	address	suggest	baseline	optimize	outperform
show	design	explore	analyze	achieve	maximize	efficient	generalize
imply	reduce	propose	explain	perform	minimize	effective	understand
study	metric	observe	benefit	improve	increase	introduce	investigate
bound	better	present	compare	dataset	decrease	interpret	demonstrate
apply	result	develop	measure	evaluate	discover	experiment	state-of-the-art

Table 11: Predefined keywords for filtering sentences.

The first attempts we made to directly adjust BART for long text either expanded its positional encodings or segmented the input text and dealt with each segment individually. Below are three ways we attempted.

Arc-I: Position Encoding Expanded BART Since the original BART model is pre-trained on 512 sequence length and fintuned on 1024 sequence length.²⁸ We followed this approach and tried copying the first 1024 position encodings periodically for longer sequence and finetuned the model on our own dataset.

Arc-II: Independently-windowed BART In this architecture, we simply chunked the documents into multiple windows with 1024 window size, and then use BART to encode them separately. The final output of the encoder side is the concatenation of those window outputs. The decoder can then generate texts as normal while attending to the whole input representations.

Arc-III: Dependently-windowed BART In **Arc-II**, we ignore the interdependence between each chunk which may lead to incoherence in generated texts. Here, to model the inter-window dependencies, we followed the approach introduced in Rae et al.’s (2020). We kept a compressive memory of the past and used this memory to compute the representation of new window. The final output of the encoder side is the concatenation of those window outputs as in **Arc-II**.

However, we found that none of these adjustments could generate satisfying fluent and coherent texts according to our experiments. Common problems include interchanges

28. <https://github.com/pytorch/fairseq/issues/1413>

between first and third person narration (They... Our model...), contradiction between consecutive sentences, more descriptive texts and fewer opinions, etc.

A.9 CE Extraction Details

The basic sentence statistics of our **ASAP-Review** dataset are listed in Tab. 12.

	ICLR	NeurIPS	Both
Avg. Sentence Num.	216	198	209

Table 12: Sentence statistics of **ASAPReview** dataset. “Avg. Sentence Num.” denotes average sentence number in a paper.

We use two steps to extract salient sentences from a source document: (i) Keywords filtering, (ii) Cross-entropy method

A.9.1 KEYWORDS FILTERING

We have predefined 48 keywords and in the first stage, we select sentences containing those keywords as well as their inflections. The 48 keywords are shown in Tab. 11. After applying keywords filtering, the statistics of selected sentences are shown in Tab. 13.

	ICLR	NeurIPS	Both
Avg. Sentence Num.	97	85	92

Table 13: Sentence statistics of selected sentences after keywords filtering. “Avg. Sentence Num.” denotes average selected sentence number in a paper.

A.9.2 CROSS ENTROPY METHOD

Following Feigenblat et al.’s (2017)’s approach in unsupervised summaization. We formalize the sentence extraction problem as a combinatorial optimization problem. Specifically, we define the performance function R as below.

$$R(S) = - \sum_{w \in S} p_S(w) \log p_S(w) \quad (8)$$

$$p_S(w) = \frac{\text{Count}(w)}{\text{Len}(S)} \quad (9)$$

Where S represents the concatenation of selected sentences, $\text{Len}(S)$ represents the number of words in S while $\text{Count}(w)$ represents the number of times w appears in S . The intuition behind this performance function is that we want to select sentences that can cover more diverse words. Note that when calculating $R(S)$, we do preprocessing steps (i.e. lowercasing, removing punctuation, removing stop words etc.).

For each paper containing n sentences, we aim to find a binary vector $p = (p_1, \dots, p_n)$ in which p_i indicates whether the i -th sentence is selected such that the combination of selected

sentences achieves highest performance score and also contains fewer than 30^{29} sentences. We did this by using Cross Entropy Method (Rubinstein & Kroese, 2013). The algorithm is shown below.

1. For each paper containing n sentences, we first assume that each sentence is equally likely to be selected. We start with $p_0 = (1/2, 1/2, \dots, 1/2)$. Let $t := 1$.
2. Draw a sample X_1, \dots, X_N of Bernoulli vectors with success probability vector p_{t-1} . For each vector, concatenate the sentences selected and get N sequences S_1, \dots, S_N . Calculate the performance scores $R(S_i)$ for all i , and order them from smallest to biggest, $R_{(1)} \leq R_{(2)} \leq \dots \leq R_{(N)}$. Let γ_t be $(1-\rho)$ sample quantile of the performances: $\gamma_t = R_{(\lceil (1-\rho)N \rceil)}$.
3. Use the same sample to calculate $\hat{p}_t = (\hat{p}_{t,1}, \dots, \hat{p}_{t,n})$ via

$$\hat{p}_{t,j} = \frac{\sum_{i=1}^N I_{\{R(S_i) \geq \gamma_t\}} I_{\{X_{ij}=1\}}}{\sum_{i=1}^N I_{\{R(S_i) \geq \gamma_t\}}} \quad (10)$$

where $I_{\{c\}}$ takes the value 1 if c is satisfied, otherwise 0.

4. Perform a smoothed update.

$$p_t = \alpha \hat{p}_t + (1 - \alpha) p_{t-1} \quad (11)$$

5. If the value of γ_t hasn't changed for 3 iterations, then stop. Otherwise, set $t := t + 1$ and return to step 2.

The elements in p_t will converge to either very close to 0 or very close to 1. And we can sample from the converged p_t to get our extraction.

We chose $N = 1000$, $\rho = 0.05$ and $\alpha = 0.7$ when we ran this algorithm. If we happen to select more than 30 sentences in a sample, we drop this sample. Note that we slightly decrease the initial probability when there are more than 90 sentences after filtering to ensure enough sample number in the first few iterations.

A.10 Detailed Analysis and Case Study

We take our aspect-enhanced model using CE extraction to conduct case study. Tab. 16 lists five examples for each aspect the model mentions. It can be seen that the language use of generated reviews are pretty close to real reviewers.

Evidence-sensitive For aspect-enhanced model, It would also be interesting to trace back to the evidence when the model generates a specific aspect. To do that we inspect where the model attends when it generates a specific aspect by looking at the attention values with respect to the source input.³⁰

29. This number is chosen according to our empirical observations. We need to extract sentences that can fit BART's input length (1024).
 30. The way we aggregate attention values is to take the maximum, no matter is to aggregate tokens to a word or to aggregate different attention heads or to aggregate words to an aspect.

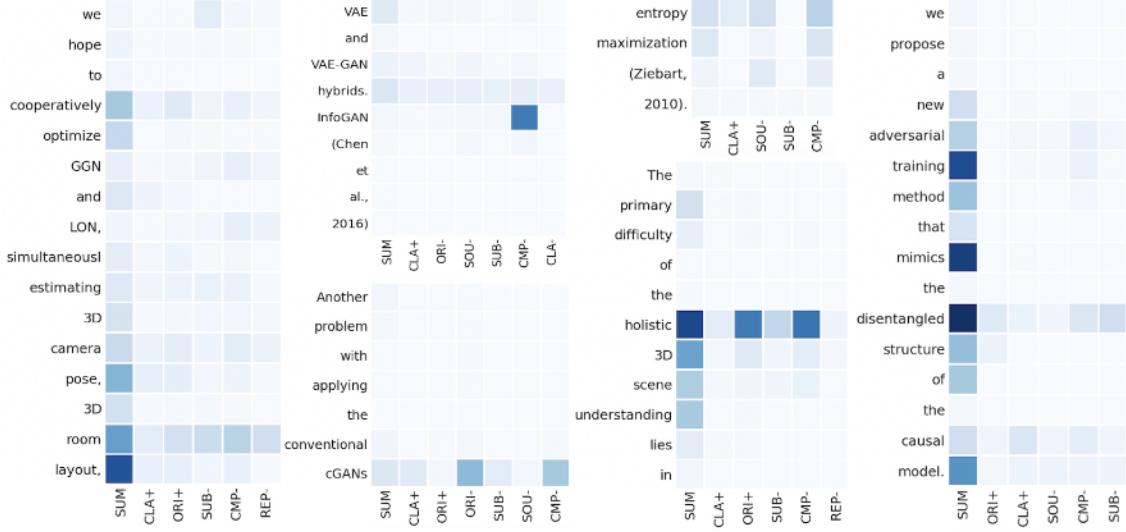


Figure 9: Attention heatmap between source document and generated reviews. + denotes positive sentiment and – denotes negative sentiment.

And interestingly, we found that the model attends to the reasonable place when it generates a specific aspect. Fig. 9 presents the attention heatmap of several segment texts, the bottom of the figure shows aspects the model generates. There are some common patterns we found when we examined the attention values between the source input and output.

1. When the model generates summary, it will attend to sentences that contain strong indicators like “we propose” or “we introduce”.
2. When it generates originality, it will attend to previous work part as well as places describing contributions of this work.
3. When it generates substance, it will attend to experiment settings and number of experiments conducted.
4. When it generates meaningful comparison, it will attend to places contains “et al.”.

A.11 Calculation of Aspect Score

For accepted (rejected) papers, we calculate the average aspect score for each aspect.

The aspect score of a review is calculated as follows.

- If an aspect does not appear in a review, then we count the score for this aspect as 0.5 (which stands for neutral).
- If an aspect appears in a review, we denote its occurrences as $\mathcal{O} = \{o_1, o_2, \dots, o_n\}$ where n is the total number of occurrences. And we denote the positive occurrences

of this aspect as $\mathcal{O}_p = \{o_{p_1}, o_{p_2}, \dots, o_{p_n}\}$ where p_n is the total number of positive occurrences. The aspect score is calculated using Formula 12.

$$\text{Aspect Score} = \frac{p_n}{n} \quad (12)$$

A.12 Disparity Analysis for All Models

Here, following the methods we proposed in §5.3.1, we list the disparity analysis for all models in Fig. 10, Fig. 11, Tab. 14, Tab. 15.

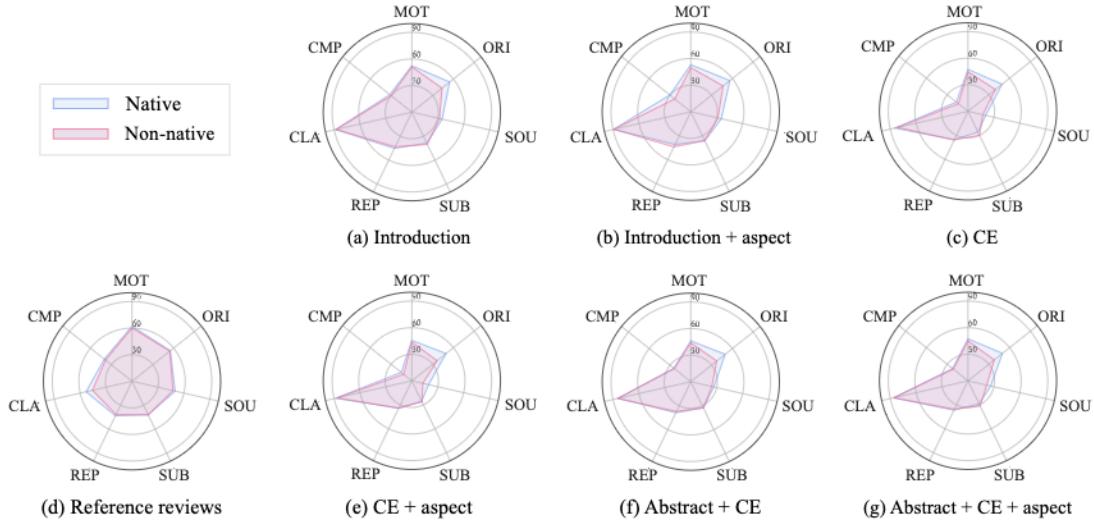


Figure 10: Spider chart of aspect scores for all models with regard to nativeness.

	MOT	ORI	SOU	SUB	REP	CLA	CMP	Total
INTRO	-0.72	+18.71	+3.84	-3.66	+0.73	-13.32	+2.40	43.39
INTRO+ASPECT	+3.12	+15.75	+6.14	+0.66	-10.61	-13.50	+19.05	68.84
CE	+2.56	+18.33	+11.16	-13.41	-3.71	-9.94	+13.49	72.58
CE+ASPECT	+1.13	+24.77	+28.78	-2.92	-3.18	-12.02	+18.36	91.18
ABSTRACT+CE	+1.77	+23.01	+3.79	+0.44	+0.37	-15.18	-2.13	46.69
ABSTRACT+CE+ASPECT	+1.72	+22.23	+12.94	-8.30	-0.38	-13.40	+0.89	59.86

Table 14: Disparity differences regarding nativeness. Total is the sum of absolute value of disparity difference.

Appendix B. Supplemental Material

B.1 Dataset Annotation Guideline

The annotation guideline for annotating aspects in reviews can be found at <https://github.com/neulab/ReviewAdvisor/blob/main/materials/AnnotationGuideline.pdf>

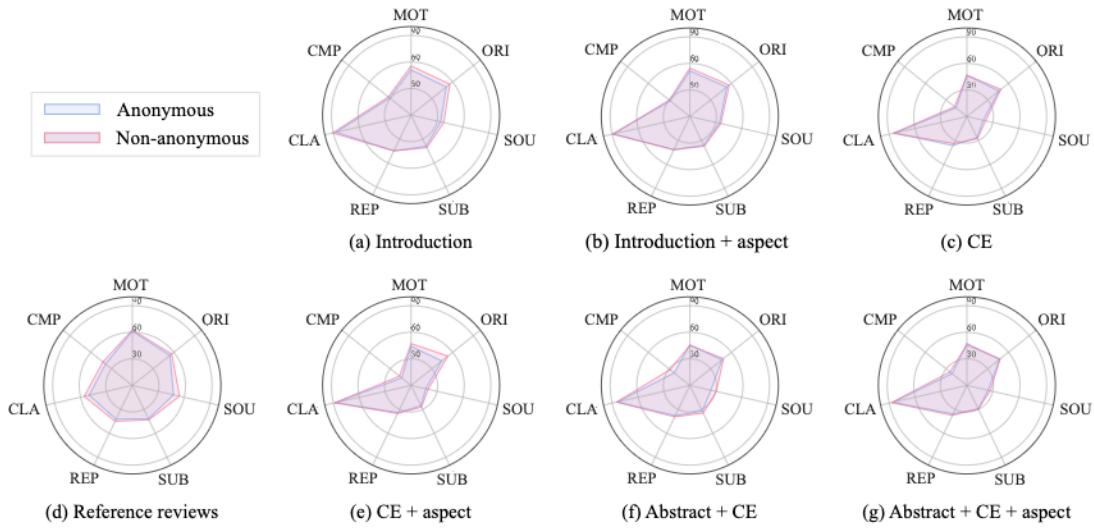


Figure 11: Spider chart of aspect scores for all models with regard to anonymity.

	MOT	ORI	SOU	SUB	REP	CLA	CMP	Total
INTRO	-5.69	-4.43	+2.76	-0.64	+5.65	+5.80	+3.02	28.00
INTRO + ASPECT	-3.53	-1.65	+7.85	+0.01	+5.93	+11.02	+4.20	34.20
CE	+1.89	-1.18	+0.05	-0.44	+13.09	+8.00	-2.56	27.21
CE+ASPECT	-4.20	-12.32	-0.52	-2.57	+2.70	+8.75	-10.31	41.37
ABSTRACT+CE	+3.18	-0.05	-7.96	-3.73	+2.25	+8.69	-12.02	37.88
ABSTRACT+CE+ASPECT	+5.45	+2.49	+2.80	+5.69	+1.33	+8.03	-3.79	29.59

Table 15: Disparity differences regarding anonymity. Total is the sum of absolute value of disparity difference.

Motivation

1. The motivation of using the conditional prior is unclear.
 2. I think this paper will be of interest to the NeurIPS community.
 3. The idea of continual learning is interesting and the method is well motivated.
 4. Overall, I think this paper is a good contribution to the field of adversarial robustness.
 5. It is hard to understand the motivation of the paper and the motivation behind the proposed methods.
-

Originality

1. This paper presents a novel approach to cross-lingual language model learning.
 2. The novelty of the paper is limited . The idea of using low rank matrices is not new.
 3. The proposed method seems to be very similar to the method of Dong et al. (2018).
 4. The idea of using neural networks to learn edit representations is interesting and novel .
 5. The proposed method seems to be a simple extension of the batched-E-step method proposed by Shazeer et al.
-

Soundness

1. This assumption is not true in practice .
 2. The experimental results are not very convincing .
 3. But the authors do not provide any theoretical justification for this claim.
 4. The theoretical results are sound and the experimental results are convincing.
 5. The paper does not provide any insights on the reasons for the success of the supervised methods.
-

Substance

1. The experiments are well-conducted.
 2. The ablation study in Section A.1.1 is not sufficient.
 3. It would be better to show the performance on a larger dataset.
 4. The authors should show the performance on more difficult problems.
 5. The experiments are extensive and show the effectiveness of the proposed method.
-

Replicability

1. It is not clear how the network is trained.
 2. The authors should provide more details about the experiments.
 3. The authors should provide more details about the hyperparameters.
 4. The authors should provide more details about the training procedure.
 5. It would be better if the authors can provide more details about the hyperparameters of LST.
-

Meaningful Comparison

1. The author should compare with [1 , 2 , 3] and [4].
 2. The authors should compare the proposed method with existing methods .
 3. It would be more convincing if the authors can compare with other methods such as AdaGrad.
 4. Authors should compare the method with the state-of-the-art methods in real-world applications .
 5. I also think the paper should compare the performance of intrinsic fear with the other methods proposed in [1 , 2 , 3 , 4 , 5].
-

Clarity

1. There are some typos in the paper.
 2. The paper is well-written and easy to follow.
 3. It is not clear to me how to interpret the results in Table 1.
 4. It would be better if the authors can provide a more detailed explanation of the difference.
 5. The paper is not well organized . It is hard to follow the description of the proposed method.
-

Table 16: Examples for different aspect mention from generated reviews.

References

- Angelidis, S., & Lapata, M. (2018). Summarizing opinions: Aspect extraction meets sentiment prediction and they are both weakly supervised. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 3675–3686, Brussels, Belgium. Association for Computational Linguistics.
- Anjum, O., Gong, H., Bhat, S., Hwu, W.-M., & Xiong, J. (2019). PaRe: A paper-reviewer matching approach using a common topic space. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 518–528, Hong Kong, China. Association for Computational Linguistics.
- August, T., Kim, L., Reinecke, K., & Smith, N. A. (2020). Writing strategies for science communication: Data and computational analysis. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 5327–5344, Online. Association for Computational Linguistics.
- Bartoli, A., De Lorenzo, A., Medvet, E., & Tarlao, F. (2016). Your paper has been accepted, rejected, or whatever: Automatic generation of scientific paper reviews. In *International Conference on Availability, Reliability, and Security*, pp. 19–28. Springer.
- Beltagy, I., Lo, K., & Cohan, A. (2019). Scibert: A pretrained language model for scientific text. *arXiv preprint arXiv:1903.10676*.
- Bolukbasi, T., Chang, K.-W., Zou, J., Saligrama, V., & Kalai, A. (2016). Man is to computer programmer as woman is to homemaker? debiasing word embeddings..
- Bornmann, L., & Mutz, R. (2015). Growth rates of modern science: A bibliometric analysis based on the number of publications and cited references. *Journal of the Association for Information Science and Technology*, 66(11), 2215–2222.
- Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. *arXiv preprint arXiv:2005.14165*.
- Cachola, I., Lo, K., Cohan, A., & Weld, D. (2020a). TLDR: Extreme summarization of scientific documents. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 4766–4777, Online. Association for Computational Linguistics.
- Cachola, I., Lo, K., Cohan, A., & Weld, D. S. (2020b). Tldr: Extreme summarization of scientific documents. *ArXiv, abs/2004.15011*.
- Chakraborty, S., Goyal, P., & Mukherjee, A. (2020). Aspect-based sentiment analysis of scientific reviews. *arXiv preprint arXiv:2006.03257*.
- Chen, Y.-C., & Bansal, M. (2018). Fast abstractive summarization with reinforce-selected sentence rewriting. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Vol. 1, pp. 675–686.
- Cohan, A., Dernoncourt, F., Kim, D. S., Bui, T., Kim, S., Chang, W., & Goharian, N. (2018a). A discourse-aware attention model for abstractive summarization of long documents. In *NAACL-HLT*.

- Cohan, A., Dernoncourt, F., Kim, D. S., Bui, T., Kim, S., Chang, W., & Goharian, N. (2018b). A discourse-aware attention model for abstractive summarization of long documents. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pp. 615–621, New Orleans, Louisiana. Association for Computational Linguistics.
- Cohan, A., & Goharian, N. (2017). Scientific article summarization using citation-context and article's discourse structure. *arXiv preprint arXiv:1704.06619*.
- De Bellis, N. (2009). *Bibliometrics and citation analysis: from the science citation index to cybermetrics*. scarecrow press.
- Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 4171–4186.
- Dou, Z.-Y., Liu, P., Hayashi, H., Jiang, Z., & Neubig, G. (2020). Gsum: A general framework for guided neural abstractive summarization. *arXiv preprint arXiv:2010.08014*.
- Efron, B. (1992). Bootstrap methods: another look at the jackknife. In *Breakthroughs in statistics*, pp. 569–593. Springer.
- Erera, S., Shmueli-Scheuer, M., Feigenblat, G., Nakash, O., Boni, O., Roitman, H., Cohen, D., Weiner, B., Mass, Y., Rivlin, O., Lev, G., Jerbi, A., Herzig, J., Hou, Y., Jochim, C., Gleize, M., Bonin, F., & Konopnicki, D. (2019). A summarization system for scientific documents. In *EMNLP/IJCNLP*.
- Feigenblat, G., Roitman, H., Boni, O., & Konopnicki, D. (2017). Unsupervised query-focused multi-document summarization using the cross entropy method. In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR '17*, p. 961–964, New York, NY, USA. Association for Computing Machinery.
- Frermann, L., & Klementiev, A. (2019). Inducing document structure for aspect-based summarization. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 6263–6273, Florence, Italy. Association for Computational Linguistics.
- Gao, Y., Eger, S., Kuznetsov, I., Gurevych, I., & Miyao, Y. (2019). Does my rebuttal matter? insights from a major NLP conference. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 1274–1290, Minneapolis, Minnesota. Association for Computational Linguistics.
- Gehrmann, S., Deng, Y., & Rush, A. (2018). Bottom-up abstractive summarization. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 4098–4109.
- Guu, K., Lee, K., Tung, Z., Pasupat, P., & Chang, M.-W. (2020). Realm: Retrieval-augmented language model pre-training. *arXiv preprint arXiv:2002.08909*.

- Hayashi, H., Budania, P., Wang, P., Ackerson, C., Neervannan, R., & Neubig, G. (2020). Wikiasp: A dataset for multi-domain aspect-based summarization. *Transactions of the Association for Computational Linguistics (TACL)*.
- He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778.
- Hermann, K. M., Kočiský, T., Grefenstette, E., Espeholt, L., Kay, W., Suleyman, M., & Blunsom, P. (2015). Teaching machines to read and comprehend. In *Advances in Neural Information Processing Systems*, pp. 1684–1692.
- Hou, Y., Jochim, C., Gleize, M., Bonin, F., & Ganguly, D. (2019). Identification of tasks, datasets, evaluation metrics, and numeric scores for scientific leaderboards construction. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 5203–5213, Florence, Italy. Association for Computational Linguistics.
- Huang, J.-B. (2018). Deep paper gestalt. *arXiv preprint arXiv:1812.08775*.
- Jain, S., van Zuylen, M., Hajishirzi, H., & Beltagy, I. (2020). SciREX: A challenge dataset for document-level information extraction. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 7506–7516, Online. Association for Computational Linguistics.
- Jecmen, S., Zhang, H., Liu, R., Shah, N. B., Conitzer, V., & Fang, F. (2020). Mitigating manipulation in peer review via randomized reviewer assignments. *arXiv preprint arXiv:2006.16437*.
- Jefferson, T., Alderson, P., Wager, E., & Davidoff, F. (2002a). Effects of editorial peer review: a systematic review. *Jama*, 287(21), 2784–2786.
- Jefferson, T., Wager, E., & Davidoff, F. (2002b). Measuring the quality of editorial peer review. *Jama*, 287(21), 2786–2790.
- Jha, R., Abu-Jbara, A., & Radev, D. (2013). A system for summarizing scientific topics starting from keywords. In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pp. 572–577, Sofia, Bulgaria. Association for Computational Linguistics.
- Jha, R., Coke, R., & Radev, D. R. (2015a). Surveyor: A system for generating coherent survey articles for scientific topics. In Bonet, B., & Koenig, S. (Eds.), *Proceedings of the Twenty-Ninth AAAI Conference on Artificial Intelligence, January 25-30, 2015, Austin, Texas, USA*, pp. 2167–2173. AAAI Press.
- Jha, R., Finegan-Dollak, C., King, B., Coke, R., & Radev, D. (2015b). Content models for survey generation: A factoid-based evaluation. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 441–450, Beijing, China. Association for Computational Linguistics.
- Jin, J., Geng, Q., Zhao, Q., & Zhang, L. (2017). Integrating the trend of research interest for reviewer assignment. In *Proceedings of the 26th International Conference on World Wide Web Companion*, pp. 1233–1241.

- Kang, D., Ammar, W., Dalvi, B., van Zuylen, M., Kohlmeier, S., Hovy, E., & Schwartz, R. (2018). A dataset of peer reviews (peerread): Collection, insights and nlp applications. In *Meeting of the North American Chapter of the Association for Computational Linguistics (NAACL)*, New Orleans, USA.
- Kingma, D., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.
- Koncel-Kedziorski, R., Bekal, D., Luan, Y., Lapata, M., & Hajishirzi, H. (2019). Text Generation from Knowledge Graphs with Graph Transformers. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pp. 2284–2293, Minneapolis, Minnesota. Association for Computational Linguistics.
- Langford, J., & Guzdial, M. (2015). The arbitrariness of reviews, and advice for school administrators..
- Lewis, M., Liu, Y., Goyal, N., Ghazvininejad, M., Mohamed, A., Levy, O., Stoyanov, V., & Zettlemoyer, L. (2019). Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. *ArXiv, abs/1910.13461*.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Kütter, H., Lewis, M., Yih, W.-t., Rocktäschel, T., et al. (2020). Retrieval-augmented generation for knowledge-intensive nlp tasks. *arXiv preprint arXiv:2005.11401*.
- Lin, C.-Y., & Hovy, E. (2003). Automatic evaluation of summaries using n-gram co-occurrence statistics. In *Proceedings of the 2003 Human Language Technology Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 150–157.
- Lo, K., Wang, L. L., Neumann, M., Kinney, R., & Weld, D. (2020). S2ORC: The semantic scholar open research corpus. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 4969–4983, Online. Association for Computational Linguistics.
- Luan, Y., He, L., Ostendorf, M., & Hajishirzi, H. (2018). Multi-task identification of entities, relations, and coreference for scientific knowledge graph construction. *arXiv preprint arXiv:1808.09602*.
- Luu, K., Koncel-Kedziorski, R., Lo, K., Cachola, I., & Smith, N. A. (2020). Citation text generation. *ArXiv, abs/2002.00317*.
- Manzoor, E., & Shah, N. B. (2020). Uncovering latent biases in text: Method and application to peer review..
- Mohammad, S., Dorr, B., Egan, M., Hassan, A., Muthukrishnan, P., Qazvinian, V., Radev, D., & Zajic, D. (2009). Using citations to generate surveys of scientific paradigms. In *Proceedings of Human Language Technologies: The 2009 Annual Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 584–592, Boulder, Colorado. Association for Computational Linguistics.
- Nallapati, R., Zhai, F., & Zhou, B. (2017). Summarunner: A recurrent neural network based sequence model for extractive summarization of documents. *ArXiv, abs/1611.04230*.

- Narayan, S., Cohen, S. B., & Lapata, M. (2018). Don't give me the details, just the summary! Topic-aware convolutional neural networks for extreme summarization. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, Brussels, Belgium.
- Nguyen, J., Sánchez-Hernández, G., Agell, N., Rovira, X., & Angulo, C. (2018). A decision support tool using order weighted averaging for conference review assignment. *Pattern Recognition Letters*, 105, 114–120.
- Paulus, R., Xiong, C., & Socher, R. (2017). A deep reinforced model for abstractive summarization. *arXiv preprint arXiv:1705.04304*.
- Qiao, F., Xu, L., & Han, X. (2018). Modularized and attention-based recurrent convolutional neural network for automatic academic paper aspect scoring. In *International Conference on Web Information Systems and Applications*, pp. 68–76. Springer.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language models are unsupervised multitask learners. *OpenAI blog*, 1(8), 9.
- Rae, J. W., Potapenko, A., Jayakumar, S. M., & Lillicrap, T. (2020). Compressive transformers for long-range sequence modelling. *ArXiv*, abs/1911.05507.
- Rogers, A., & Augenstein, I. (2020). What can we do to improve peer review in NLP?. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 1256–1262, Online. Association for Computational Linguistics.
- Rubinstein, R. Y., & Kroese, D. P. (2013). *The cross-entropy method: a unified approach to combinatorial optimization, Monte-Carlo simulation and machine learning*. Springer Science & Business Media.
- Smith, R. (2006). Peer review: A flawed process at the heart of science and journals. *Journal of the Royal Society of Medicine*, 99, 178 – 182.
- Stanovsky, G., Smith, N. A., & Zettlemoyer, L. (2019). Evaluating gender bias in machine translation. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 1679–1684, Florence, Italy. Association for Computational Linguistics.
- Stelmakh, I., Shah, N., & Singh, A. (2019). On testing for biases in peer review. In *Advances in Neural Information Processing Systems*, pp. 5286–5296.
- Subramanian, S., Li, R., Pilault, J., & Pal, C. (2019). On extractive and abstractive neural document summarization with transformer language models. *arXiv preprint arXiv:1909.03186*.
- Tabah, A. N. (1999). Literature dynamics: Studies on growth, diffusion, and epidemics. *Annual review of information science and technology (ARIST)*, 34, 249–86.
- Tomkins, A., Zhang, M., & Heavlin, W. D. (2017). Reviewer bias in single-versus double-blind peer review. *Proceedings of the National Academy of Sciences*, 114(48), 12708–12713.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention is all you need. In *Advances in neural information processing systems*, pp. 5998–6008.

- Von Bearnensquash, C. (2010). Paper gestalt. *Secret Proceedings of Computer Vision and Pattern Recognition (CVPR)*.
- Wadden, D., Lin, S., Lo, K., Wang, L. L., van Zuylen, M., Cohan, A., & Hajishirzi, H. (2020). Fact or fiction: Verifying scientific claims. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 7534–7550, Online. Association for Computational Linguistics.
- Wang, K., & Wan, X. (2018). Sentiment analysis of peer review texts for scholarly papers. In *The 41st International ACM SIGIR Conference on Research & Development in Information Retrieval*, pp. 175–184.
- Wang, Q., Zeng, Q., Huang, L., Knight, K., Ji, H., & Rajani, N. F. (2020a). Reviewrobot: Explainable paper review generation based on knowledge synthesis. In *Proceedings of INLG*.
- Wang, Q., Zeng, Q., Huang, L., Knight, K., Ji, H., & Rajani, N. F. (2020b). ReviewRobot: Explainable paper review generation based on knowledge synthesis. In *Proceedings of the 13th International Conference on Natural Language Generation*, pp. 384–397, Dublin, Ireland. Association for Computational Linguistics.
- Xiao, W., & Carenini, G. (2019). Extractive summarization of long documents by combining global and local context. *ArXiv, abs/1909.08089*.
- Xing, X., Fan, X., & Wan, X. (2020). Automatic generation of citation texts in scholarly papers: A pilot study. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 6181–6190, Online. Association for Computational Linguistics.
- Xiong, W., & Litman, D. (2011). Automatically predicting peer-review helpfulness. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pp. 502–507, Portland, Oregon, USA. Association for Computational Linguistics.
- Yasunaga, M., Kasai, J., Zhang, R., Fabbri, A. R., Li, I., Friedman, D., & Radev, D. R. (2019a). Scisummnet: A large annotated corpus and content-impact models for scientific paper summarization with citation networks. In *AAAI*.
- Yasunaga, M., Kasai, J., Zhang, R., Fabbri, A. R., Li, I., Friedman, D., & Radev, D. R. (2019b). Scisummnet: A large annotated corpus and content-impact models for scientific paper summarization with citation networks. In *The Thirty-Third AAAI Conference on Artificial Intelligence, AAAI 2019, The Thirty-First Innovative Applications of Artificial Intelligence Conference, IAAI 2019, The Ninth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2019, Honolulu, Hawaii, USA, January 27 - February 1, 2019*, pp. 7386–7393. AAAI Press.
- Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., & Artzi, Y. (2019). Bertscore: Evaluating text generation with bert. *arXiv, arXiv-1904*.
- Zhao, J., Wang, T., Yatskar, M., Ordonez, V., & Chang, K.-W. (2018). Gender bias in coreference resolution: Evaluation and debiasing methods. *arXiv preprint arXiv:1804.06876*.